

High-Performance
Computing Center
Stuttgart

MPI - Introduction

EuroMPI/USA 2025 Tutorial
Christoph Niethammer & Joseph Schuchart


Overview

H L R **I** S

- Motivation
- History of the MPI Interface
- MPI Basics
- Point to Point (P2P) communication
- Collective operations
- MPI and GPU
- Process Topologies
- MPI + Threads & Partitioned communication
- MPI one-sided communication / RMA

Motivation

- Supercomputer built as distributed memory systems with $O(1000)$ nodes
- Many different HPC hardware vendors
- Software portability



History of the MPI interface

- 1994: meeting of 40 universities and companies: Standardization MPI-1.0 based on experiences of existing communication APIs → Foundation of the **MPI-Forum**
- 1995: MPI-1.1 minor corrections
- 1997: MPI-1.2 corrections, rationals, minor corrections
- 1997: MPI-2.0 large body of changes: MPI-IO, one-sided, dynamic process management, ...
- 2008: MPI-1.3: ammendments and cleanup in MPI-1
- 2009: MPI-2.1 lots of corrections on MPI-2
- 2009: MPI-2.2 corrections and simplifications including deprecation of the C++ Interface
- 2012: MPI-3.0 many new features: non-blocking collectives, RMA (aka true one-sided), new Fortran Interface
- 2015: MPI-3.1 minor corrections; few new routines
- 2021: MPI-4.0 Partitioned Communication, „Big-Count“
- 2023: MPI 4.1 corrections
- 2025: MPI 5.0 MPI ABI
- 2025+: MPI next ... MPI Forum just started working on it :)

MPI Implementations

The MPI Standard defines the syntax and semantics of communication functions but not the actual implementation!

There are different ways implementing MPI:

- **MPICH:**

from Argonne National Labs (ANL)
the first available implementation
foundation of many MPIs

MVAPICH2

NEC's MPI-SX

Intel MPI

HPE/Cray's MPI

IBM's MPI (BG/L, discontinued)

- **Open MPI:**

from Labs, Industry & Universities
the „new“ implementation



Sun's (Oracle's) HPC Cluster Tools

Fujitsu's MPI for K-Computer

Bull MPI

IBM Spectrum MPI

Nvidia HPC-X

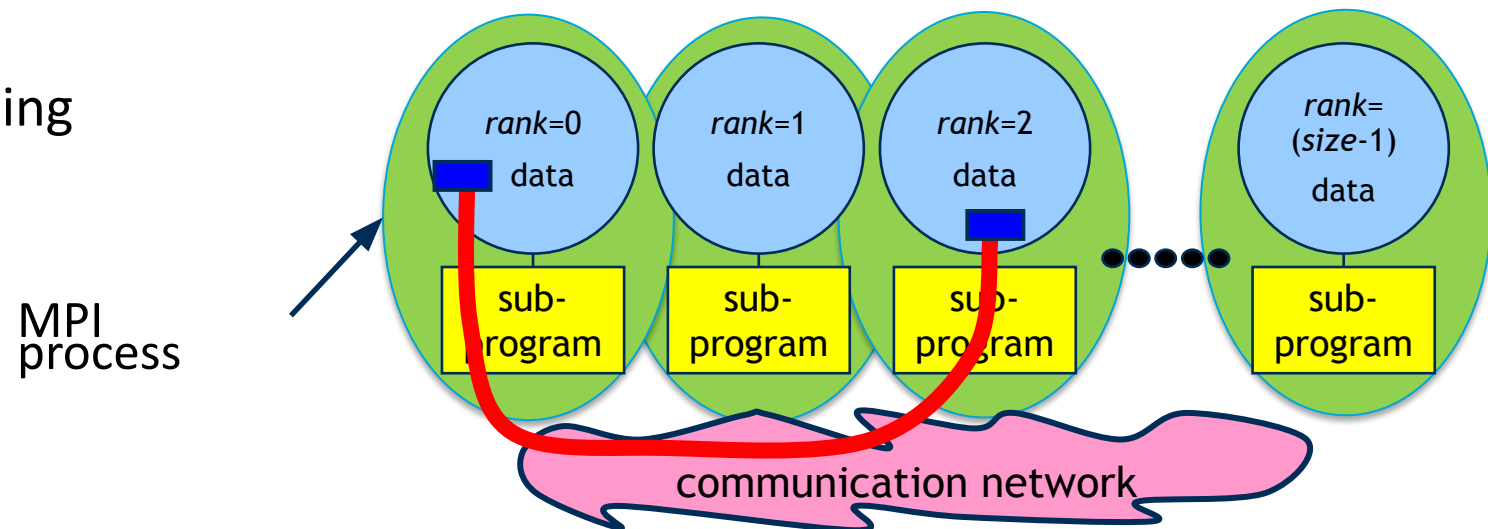
Note: MPI includes an ABI specification since MPI 5.0 allowing interoperability between implementations

Message Passing Interface (MPI)

Basics

MPI Basics

- Each processor in a message passing program runs a sub-program referred to as **MPI process**
 - Written in a conventional sequential language, e.g., C or Fortran,
 - Typically the same on each processor (SPMD) but can also be different (MPMD)
- All communication, work and data distribution is based on value of **rank** that identifies a MPI process in a **Group of MPI processes**
- Communication based on exchanging messages between MPI processes
→ message passing



Functionalities of the MPI

An abundance of communication functionality:

- Powerful functions to **group and structure MPI processes**, e.g. for **topologies** or structural close-ness of underlying physical problem at hand!
- **Blocking Point-to-Point (P2P)**, e.g. MPI_Send/MPI_Recv:
Return to caller, as soon as the buffer may be reused by the caller!
- **Non-Blocking (Immediate) P2P**, e.g. MPI_Isend/MPI_Irecv:
Communication „in background“, for overlapping of computation & communication!
- **Collective Communications**, e.g. MPI_Bcast, MPI_Reduce, but also MPI_Allgatherv, i.e. „All“ means all processes of so-called communicator participate, „gather“ means data is gathered from all processes and „v“ means each process gathers variable amount of data!
- **Remote-Memory-Access (RMA)**: Direct memory read, or write or even atomic increment of remote memory; with the proper Hardware-support this is way faster than P2P Communication...
- **Parallel File-IO**, e.g. MPI_File_open, MPI_File_read_all...

Simple MPI program example

H L R **I** S

```
#include "mpi.h"
int main (int argc, char * argv[]) {
    int rank, size, sndbuf, rcvbuf;
    MPI_Comm comm;
    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_dup(MPI_COMM_WORLD, &comm);
    MPI_Comm_rank(comm, &rank);
    MPI_Comm_size(comm, &size);
    if (0 == rank) {
        sndbuf=42;
        MPI_Send(&sndbuf, 1, MPI_INT, 1, 4711, comm);
    } else if (1 == rank) {
        MPI_Recv(&rcvbuf, 1, MPI_INT, 0, 4711, comm, &status);
    }
    MPI_Finalize ();
}
```

Initialize MPI with the
„**world programming
model**“

Create a communication
context for a group of MPI
processes, i.e., **communicator**

Administrative functions

Communication, here P2P
between MPI processes with
rank 0 and 1 in the group of
comm

MPI program generation

- The compiler's knowledge about the MPI API comes from the interface definitions in the header file (#include 'mpi.h') in C or the module (use mpi_f08) in Fortran
- At the end the linker has to link in the MPI library

```
gcc -I <MPI_INSTALLATION_DIR>/include \  
    -o mpi_prog mpi_prog.c \  
    -L <MPI_INSTALLATION_DIR>/lib -lmpi
```

↑
Path to
include directory

↑
Path to
library directory

↑
The library to be linked to, here:
libmpi.so

MPI libraries come with compiler wrappers, e.g., `mpicc` as replacements for the original compiler commands to simplify this

```
mpicc -o mpi_prog mpi_prog.c
```

Note: For applications using the cmake build system, the FindMPI module (find(MPI)) can be used.

Executing an MPI program

H L R I S


- Requires
 - Starting of MPI processes on various resources
 - Propagating connectivity information between the MPI processes
- MPI startup command:
`mpirun -n <numprocs> <additional options> <program>`

Note: MPI implementations provide many additional options to `mpirun`, e.g., for controlling process placement

MPI program execution example using Open MPI

The following command starts the program execution of `mpi_prog` using 4 processes, passing 2 arguments.

```
mpiexec -n 4 ./mpi_prog arg1 arg2
```



- Upon first call of `MPI_Init()` processes „get to know“ each other following the “world model”
When the processes return from `MPI_Init()`, `MPI_COMM_WORLD` is defined and a process's rank in it is known.
- **Where** the processes are executed, i.e., on which node / which core, is MPI implementation-dependent. All MPI implementations allow to control this *mapping*.

Availability of MPI for other programming languages

H L R **I** S

- MPI defines only a C and Fortran interface
- There exist many „Language Bindings“ – typically based on the C-Interface
 - C++: via the C-Interfaces or, Boost.MPI (MPI 1.1), MPL (MPI 3.1), ...
 - Rust: rsmapi (MPI 3.1, <https://docs.rs/mapi/latest/mapi/>)
 - Python: mapi4py (<https://mapi4py.readthedocs.io/>)
 - Java (z.B., in Open MPI via import mapi.*)
- All of those require the availability of an MPI library installation (C/Fortran)


Message Passing Interface (MPI)

Point to Point (P2P) communication

Point to Point (P2P) communicatoin


- Communication between two processes.
- Source process sends message to destination process.
- Communication takes place within a communicator, e.g., MPI_COMM_WORLD.
- Processes are identified by their ranks in the communicator.

communicator



Messages

H L R I S



- Messages are packets of data moving between sub-programs
- Necessary information for the message passing system:
 - sending process
 - source location
 - source data type
 - source data size
 - source tag
 - receiving process i.e., the ranks
 - destination location
 - destination data type
 - destination buffer size
 - recv tag



Point to Point (P2P) Communication

Simplest form of MPI communication is **blocking point-to-point**

- One process sends, another process receives: „two-sided Communication“.
- The MPI Standard defines for each call the API as:

```
int MPI_Send(void * buf, int count, MPI_Datatype ddt, int rank, int tag, MPI_Comm comm);
```

Send the buffer pointed to by `buf` with `count` elements of type `ddt` to process `rank` within the communicator `comm` identified by `tag`.

The call returns as soon as `buf` may be reused by the application !

```
int MPI_Recv(void * buf, int count, MPI_Datatype ddt, int rank, int tag, MPI_Comm comm,  
             MPI_Status * status);
```

Receive into the buffer pointed to by `buf` with max. `count` elements of type `ddt` from process `rank` within the communicator `comm` identified by `tag`.

The call returns as soon as `buf` contains the complete data!


Example: MPI Ping-Pong

H L R I S

```
#include "mpi.h"
int main (int argc, char * argv[]) {
    int rank, size, buf = 42;
    MPI_Comm comm;
    MPI_Init(&argc, &argv);
    MPI_Comm_dup(MPI_COMM_WORLD, &comm);
    MPI_Comm_rank(comm, &rank);
    MPI_Comm_size(comm, &size);
    if (0 == rank) {
        MPI_Send(&buf, 1, MPI_INT, 1, 4711, comm);
        MPI_Recv(&buf, 1, MPI_INT, 1, 4711, comm, MPI_STATUS_IGNORE);
    } else if (1 == rank) {
        MPI_Recv(&buf, 1, MPI_INT, 0, 4711, comm, MPI_STATUS_IGNORE);
        MPI_Recv(&buf, 1, MPI_INT, 0, 4711, comm, &status);
    }
    MPI_Finalize ();
}
```

rank=0

rank=1



MPI Basics – Communicator & Rank


The **rank** of each process is **defined within a group of a communicator!**

Predefined communicators are made available, e.g., in the “world model” by `MPI_Init`.

The ranks of all processes within the communicator do not change

The number of processes started with `mpirun` match exactly the range from 0 to `sizeof(MPI_COMM_WORLD) - 1`

```
mpirun -np 4 ./mpi_prog
```



Pre-defined communicators in the “world model”:

- `MPI_COMM_WORLD`: all processes
- `MPI_COMM_SELF`: The “own” process
- `MPI_COMM_NULL`: no process included

MPI communicators

- Each communication is relative to a specific **communicator** that provides the communication context!
- The predefined `MPI_COMM_WORLD` in the “world model” is a starting point used very often, ...
- There’s different methods to create communicators:
 1. `MPI_Comm_dup` duplicates the passed communicator (i.e. inside of your own library – or for usage with threads! The newly created communicator is semantically different from the old comm.
 2. Using `MPI_Comm_split_type` one may include/exclude specific processes from the newly created communicator.

```
MPI_Comm_split_type(MPI_COMM_WORLD, MPI_COMM_TYPE_SHARED, rank, &comm);  
// Splits into communicators that allow shared memory access (from MCW)
```
 3. By a way of moving processes to groups, any kind of grouping is achieved, e.g. here multiple „Servers“ and 42 „Clients“ per server.

```
MPI_Comm_group (MCW, &group_mcw); // mcw is MPI_COMM_WORLD  
for (i=0; i < size/42; i++) servers[i] = i*42;  
MPI_Group_excl (group_mcw, i, servers, group_client);  
MPI_Group_incl (group_mcw, i, servers, group_server);
```
 4. For even more flexibility groups can be obtained from MPI Sessions in the “sessions model” (see MPI 4.0)

Ranks


- The rank identifies different processes.
- The rank is the basis for any work and data distribution.

C

Fort
ran

Python

- C/C++: `int MPI_Comm_rank(MPI_Comm comm, int *rank)`
- Fortran: `MPI_COMM_RANK(comm, rank, ierror)`
`mpi_f08: TYPE(MPI_Comm) :: comm`
`INTEGER :: rank; INTEGER, OPTIONAL :: ierror`
`mpi & mpif.h: INTEGER comm, rank, ierror`
- Python: `rank = comm.Get_rank()`




CALL MPI_COMM_RANK(MPI_COMM_WORLD, myrank, ierror)

Point to Point (P2P) Communication

H L R I S

Standard blocking send & receive



When `MPI_Recv` returns, we know the message is received in full.

When `MPI_Send` returns, we **do not know**, whether the receiver has received, or even yet called `MPI_Recv`!

The message may be buffered by MPI, or be on the network
MPI could therefore block, if the message is too long!

Point to Point (P2P) communication protocols

H L R **I** S

Example: Standard blocking send & receive



```
MPI_Send (&buf, len...);    long_computation();  
long_computation();         MPI_Recv (&buf, len...);
```

„Eager Protocol“

If buffer to be send is small (ca. < 64 kB): copy the message into MPI-internal buffer, send to network and return to application

→ sender continues with `long_computation`!

„Rendezvous Protocol“

If buffer to be sent is too big (e.g. > 64 kB): send first message fragment “eagerly”, then wait until the receiver is ready for the rest, i.e., until the call of `MPI_Recv`!

Deadlock example

H L R I S

```
#include "mpi.h"
int main (int argc, char * argv[]) {
    int rank, size, sndbuf = 42, rcvbuf;
    MPI_Comm comm;
    MPI_Init(&argc, &argv);
    MPI_Comm_dup(MPI_COMM_WORLD, &comm);
    MPI_Comm_rank(comm, &rank);
    MPI_Comm_size(comm, &size);

    int to = (rank + 1) % size
    int from = (rank + size - 1) % size
    MPI_Send(&sndbuf, 1, MPI_INT, to, 4711, comm);
    MPI_Recv(&rcvbuf, 1, MPI_INT, from , 4711, comm, MPI_STATUS_IGNORE);


    MPI_Finalize ();
}
```

Deadlock!

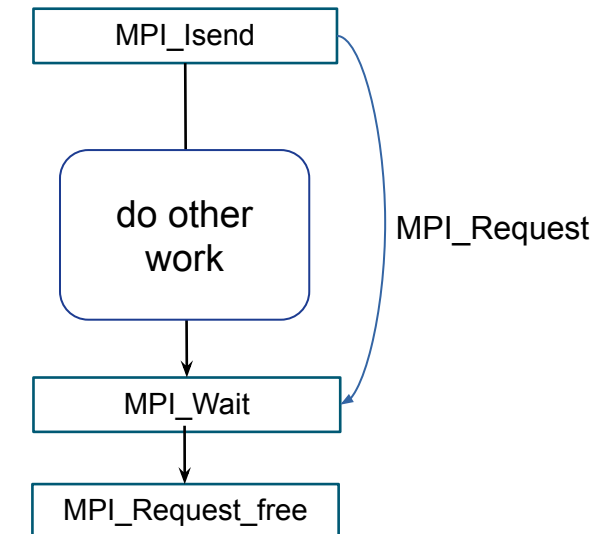
Nonblocking operations

Nonblocking operations consist of:

- A **nonblocking procedure** call: it returns immediately and allows the
- code to perform other work
- At some **later** time the sub-program must **test or wait for the completion** of the nonblocking operation



Blocking



Nonblocking

Nonblocking (immediate) P2P

```
int MPI_Isend(void * buf, int count, MPI_Datatype ddt, int rank, int tag,  
             MPI_Comm comm, MPI_Request * request);
```

Send non-blocking, i.e. `buf` with `count` elements of type `ddt` is send in “the background”!

The call returns immediately with an additional **request**.

However, **the buffer `buf` may not be touched until a corresponding `MPI_Wait()` has been called.**

Each request of a non-blocking call **has** to be finalized using either

```
MPI_Wait(), MPI_Test(),           // Wait for one request...  
MPI_Waitall(), MPI_Testall(),     // Wait for all requests to finish  
MPI_Waitsome(), MPI_Testsome(),  // Wait for possibly multiple  
MPI_Waitany(), MPI_Testany()     // Wait for one of many requests
```

Nonblocking (immediate) P2P

If you want to overlap communication & computation, to „perfectly“ hide communication, then nonblocking immediate P2P is a **must**.

An example:



```
MPI_Isend
  (&buf, 1000, MPI_INT, 1, \
   4711, MPI_COMM_WORLD, &req);
long_computation();
MPI_Wait (&req, &status);
/* Now buf may be used again */
```

```
MPI_Irecv
  (&buf, 1000, MPI_INT, 0, \
   4711, MPI_COMM_WORLD, &req);
long_computation();
MPI_Wait (&req, &status);
/* Now buf may be used again */
```

MPI Datatypes

MPI provides datatypes for message data

- Allow type transformation in inhomogeneous systems, e.g., different MPI processes run on different architectures (little/big endian)
- Writing data in platform independent file formats
- **Predefined MPI datatypes** for C and Fortran datatypes:


MPI datatype	C datatype
MPI_CHAR	char
MPI_INT	int
MPI_LONG	long
MPI_DOUBLE	double
MPI_UNSIGNED	unsigned int
...	

MPI Derived Datatypes

Out of a base data type, one may generate new derived datatype:

```
typedef struct {
    int location;
    double real;
    double imag;
} complex_loc_t cl;
```

Has the following
layout in memory



In order to send data of `complex_loc_t` type, one has to describe it using base- and derived types (here only base):

```
array_of_types[0]      = MPI_INT;
array_of_blocklen[0]   = 1;
array_of_disp[0]       = 0;
array_of_types[1]      = MPI_DOUBLE;
array_of_blocklen[1]   = 2;
array_of_disp[1]       = &(cl.real) - &(cl.location);
MPI_Type_struct(2, array_of_blocklen, array_of_disp,
               array_of_types, &complex_loc_ddt);
MPI_Type_commit(&complex_loc_ddt); // Now one may send!
```

Message Passing Interface (MPI)


Collective communication

Collective communication

- In collective communication multiple MPI processes are involved in the communication
- The group of involved processes is given by the processes in the group of the communicator
- Examples: Broadcast, Reduction, Scatter, ...
- **Synchronizing** collective communication requires that all MPI processes must start the collective MPI operation before one MPI process can finish the operation
- Examples: Allreduce, Alltoall

MPI Collective Communication - Broadcast


Broadcast data from one process to all others



```
int MPI_Bcast(void *buf, int cnt, MPI_Datatype ddt, int root, MPI_Comm comm);
```


The memory pointed to by `buf` of length `cnt*ddt` will be sent from `root` to all processes in `comm` and copied into their memory `buf`.

Attention: `MPI_Bcast()` does not have to synchronize!




Here:
cnt=4
root=0

MPI Collective Communication - Scatter




```
int MPI_Scatter(void *sendbuf, int sendcnt, MPI_Datatype send_ddt,
               void *recvbuf, int recvcnt, MPI_Datatype recv_ddt,
               int root, MPI_Comm comm);
```

Process `root` distributes to all processes the data pointed to by `sendbuf` to all processes in `comm`; the data is stored in the receiving processes in `recvbuf`.




Here:
`sendcnt=1`
`root=0`

MPI Collective Communication - Gather




```
int MPI_Gather (void *sendbuf, int sendcnt, MPI_Datatype send_ddt,
               void *recvbuf, int recvcnt, MPI_Datatype recv_ddt,
               int root, MPI_Comm comm);
```

Process `root` collects from all processes the data pointed to by `sendbuf` and saves the data in `recvbuf`.




Here:
`sendcnt=1`
`root=0`

MPI Collective Communication - Alltoall



```
int MPI_Alltoall(void *sendbuf, int sendcnt, MPI_Datatype send_ddt,
                 void *recvbuf, int recvcnt, MPI_Datatype recv_ddt,
                 MPI_Comm comm);
```


Matrix-Transposition: Each process sends `sendcnt` data for every other process within communicator, i.e. for 4 processes $4 * \text{sendcnt}$ elements.



Used for Fast Fourier Transformations (FFT)!


This is a really expensive operation!

MPI Collective Communication - Reduce



```
int MPI_Reduce(void *sendbuf, void *recvbuf, int count, MPI_Datatype ddt,
               MPI_Op op, int root, MPI_Comm comm);
```

Process `root` using the operation `op` collects and combines data from `recvbuf`. Possible ops depend on the data type, and include `MPI_MIN`, `MPI_MAX`, `MPI_SUM`, `MPI_PROD`, `MPI_LAND`, `MPI_BAND`, `MPI_LOR`, `MPI BOR`, `MPI_LXOR`, `MPI_BXOR`, `MPI_MINLOC`, `MPI_MAXLOC`. One may define one's own `MPI_Op`.



MPI Collective Communication ...

There are many variants of each of the collective calls:

- **All variant:** Instead of just `root` all processes within `comm` receive the data, e.g. `MPI_Allgather()` and `MPI_Allreduce()`
- ***v variant:** Variable number of elements per rank, i.e. each process may send a different amount of data.
- **Immediate variant:** Since MPI-3 also immediate, i.e. non-blocking variants of all calls are available, e.g. `MPI_Igather()` and `MPI_Ibarrier()`

The biggest up-point :

- **The MPI implementation may optimize the communication pattern according to the hardware, process mapping and network topology.**
- You may define your own operators

Message Passing Interface (MPI)

MPI and GPU Memory

MPI and GPUs: Implementation-Specifics

- Most MPI implementations support using GPU memory
- All MPI calls happen on the host

```
/*
 * Program that shows the use of CUDA-aware macro and runtime check.
 */
#include <stdio.h>
#include "mpi.h"

#if !defined(OPEN_MPI) || !OPEN_MPI
#error This source code uses an Open MPI-specific extension
#endif

/* Needed for MPIX_Query_cuda_support(), below */
#include "mpi-ext.h"


int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);

    printf("Compile time check:\n");
    #if defined(MPIX_CUDA_AWARE_SUPPORT) && MPIX_CUDA_AWARE_SUPPORT
        printf("This MPI library has CUDA-aware support.\n", MPIX_CUDA_AWARE_SUPPORT);
    #elif defined(MPIX_CUDA_AWARE_SUPPORT) && !MPIX_CUDA_AWARE_SUPPORT
        printf("This MPI library does not have CUDA-aware support.\n");
    #else
        printf("This MPI library cannot determine if there is CUDA-aware support.\n");
    #endif /* MPIX_CUDA_AWARE_SUPPORT */

    printf("Run time check:\n");
    #if defined(MPIX_CUDA_AWARE_SUPPORT)
        if (1 == MPIX_Query_cuda_support()) {
            printf("This MPI library has CUDA-aware support.\n");
        } else {
            printf("This MPI library does not have CUDA-aware support.\n");
        }
    #else /* !defined(MPIX_CUDA_AWARE_SUPPORT) */
        printf("This MPI library cannot determine if there is CUDA-aware support.\n");
    #endif /* MPIX_CUDA_AWARE_SUPPORT */

    MPI_Finalize();

    return 0;
}
```



MPI and GPUs: Requesting Device Support

- Sessions replace MPI_Init and can be used to request and check for device memory support
- Device memory types documented in a [side document](#)
- Info objects = string dictionaries
- Applications should request device memory kind during startup

```
int cuda_device_aware = 0;
int cuda_managed_aware = 0;
int len, flag = 0;
MPI_Info info;
MPI_Session session;
// Usage mode : REQUESTED
MPI_Info_create(&info);
MPI_Info_set(info, "mpi_memory_alloc_kinds",
               "system,cuda:device,cuda:managed");
MPI_Session_init(info, MPI_ERRORS_ARE_FATAL, &session);
MPI_Info_free(&info);
```

MPI and GPUs: Checking Device Support

- Check whether support for a specific device memory kind is **provided**
 - Application can check that **requested** memory kind is provided

```
// Usage mode : PROVIDED
MPI_Session_get_info(session, &info);
MPI_Info_get_string(info, "mpi_memory_alloc_kinds", &len, NULL, &flag);

if (flag) {
    char *val, *valptr, *kind;
    val = valptr = (char *)malloc(len);
    MPI_Info_get_string(info, "mpi_memory_alloc_kinds", &len, valptr, &flag);
    while ((kind = strsep(&val, ",")) != NULL) {
        if (strcasecmp(kind, "cuda:managed") == 0) {
            cuda_managed_aware = 1;
        } else if (strcasecmp(kind, "cuda:device") == 0) {
            cuda_device_aware = 1;
        }
    }
    free(valptr);
}
```

MPI and GPUs: Asserting Device Memory Usage

- **Assert** that only a specific device memory kind is used
 - MPI can optimize based on this assertion

```
// Usage mode : ASSERTED
MPI_Group wgroup;
MPI_Comm system_comm;
MPI_Group_from_session_pset(session, "mpi://WORLD", &wgroup);
// Create a communicator for operations on system memory
MPI_Info_create(&info);
MPI_Info_set(info, "mpi_assert_memory_alloc_kinds", "system");
MPI_Comm_create_from_group(wgroup,
    "org.mpi-forum.side-doc.mem-alloc-kind.cuda-example.system",
    info, MPI_ERRORS_ABORT, &system_comm);
MPI_Info_free(&info);
```

MPI and GPUs: Usage Example

- Allocate a CUDA device buffer and use the previously created communicator on which we asserted the use of CUDA device memory.

```
// Example: Use with OpenMP target directive
double *cuda_buffer;
cudaMalloc(sizeof(double)*N, &cuda_buffer);

#pragma omp target
{
    ...
}
MPI_Send(cuda_buffer, N, MPI_DOUBLE, peer, cuda_comm);

cudaFree(&cuda_buffer);
```

Message Passing Interface (MPI)

Process Topologies & Collective communication

MPI Process Topologies

- The linear identifier rank may not reflect the **logical communication pattern** of an application, e.g., in Cartesian process grids or graphs
- Topologies provide here a convenient naming mechanism for MPI processes in a group of processes
- Process topologies express communication path information that can help an MPI runtime in mapping MPI processes onto the underlying hardware topology
- Implemented via virtual topologies in MPI:
 - Cartesian Topology
 - Graph Topology
 - Distributed Graph Topology
- Topology information is associated to a communicator

MPI Cartesian Topologies

Cartesian Topologies represent a n-dimensional process grid

Communicator with an associated Cartesian topology is created by

```
int MPI_Cart_create(MPI_Comm comm_old, int ndims,  
    int dims[], int periods[], int reorder, MPI_Comm *comm_cart)
```

Introduces process coordinates with row-major numbering:

Hardware mapping

Example: 4 processes in 2x2 grid:

(0,0): R0	(0,1): R1
(1,0): R2	(1,1): R3

The following convenience function can factorize a number:

```
int MPI_Dims_create(int nnodes, int ndims, int dims[])
```

BUT: It takes not into account the data layout inside the application, i.e., will not optimize the amount of data to be communicated!

MPI Cartesian Topologies - coordinates

Various helpers for mapping coordinates to ranks in P2P communication functions here:

coordinate → rank:

```
int MPI_Cart_rank(MPI_Comm comm, int coords[], int *rank)
```

rank → coordinate:

```
int MPI_Cart_coords(MPI_Comm comm, int rank, int maxdims,  
                    int coords[])
```

Communication along the coordinate direction may be preformed, e.g., with
MPI_Sendrecv:

```
int MPI_Cart_shift(MPI_Comm comm, int direction,  
                  int disp, int *rank_source, int rank_dest)
```

MPI Graph Topologies

- Most general form to describe process topologies:
 - Nodes: MPI processes
 - Edges: communication
 - Edge weights: additional hints (e.g., bandwidth, latency, ...)
 - Come in two flavours:
 - Graph:
 - Each process has the full graph information
 - Distributed Graph
 - Each process only has a local subset of the graph
 - Allow specification of additional weights for edges
- ← This does not scale!

MPI Distributed Graph Topologies

Distributed graphs can be created in two ways:

- Any process can specify any part of the graph:

```
int MPI_Dist_graph_create(MPI_Comm comm_old, int n,  
    int sources[], int degrees[], int destinations[],  
    int weights[], MPI_Info info, int reorder,  
    MPI_Comm *comm_dist_graph)
```


- Each process specifies its neighbours:

```
int MPI_Dist_graph_create_adjacent(MPI_Comm comm_old,  
    int indegree, int sources[], int sourceweights[],  
    int outdegree, int destinations[], int destweights[],  
    MPI_Info info, int reorder, MPI_Comm *comm_dist_graph)
```

MPI Distributed Graph Topologies

for MPI_Dist_graph_create:

process	n	sources	degrees	destinations	weights
0	4	0,1,2,3	2,1,1,2	1,3,0,3,0,2	1,1,1,1,1,1
1	0	-	-	-	-
2	0	-	-	-	-
3	0	-	-	-	-



for MPI_Dist_graph_create_adjacent:

process	indegree	sources	sourceweights	outdegree	destinations	destweights
0	2	1,3	1,1	2	1,3	1,1
1	1	0	1	1	0	1
2	1	3	1	1	3	1
3	2	0,2	1,1	2	0,2	1,1

MPI Graph Topologies

Information about the number of neighbours for a specified rank:

```
int MPI_Graph_neighbors_count(MPI_Comm comm,  
    int rank, int *nneighbors)
```

Actual neighbours for a specified rank:

```
int MPI_Graph_neighbors(MPI_Comm comm,  
    int rank, int maxneighbors, int neighbors[])
```

MPI Neighborhood Collectives

MPI Topologies define neighbours for MPI processes.

One can use collective operations on these neighbours – and replace P2P communication.

E.g., Bcast to or Reduce over all neighbours of a MPI process instead of a loop over all neighbours and performing individual send/recvs.

Related functions are of the form **MPI_Neighbor_<collective>**

Advantage:

- Communication pattern can be optimized by the MPI libraries.
- Easier to understand and simpler code by exposing what is intended (reduce!, gather!, etc.)

Message Passing Interface (MPI)

MPI and Threads (e.g. OpenMP) + Partitioned Communication

MPI and threads

When MPI is used in combination with threads the MPI library has to protect its internal state. Therefore the MPI has to be initiated with

```
int MPI_Init_thread(int *argc, char ***argv,  
                    int required, int *provided)
```

Requested/provided can be one of

- `MPI_THREAD_SINGLE`: only one thread calls MPI functions
- `MPI_THREAD_FUNNELED`: only the thread that called `MPI_Init_thread` does MPI
- `MPI_THREAD_SERIALIZED`: only one thread at a time will call MPI functions
- `MPI_THREAD_MULTIPLE`: multiple threads may call MPI functions concurrently

Starting an hybrid MPI and threads program

- To start use still mpiexec to start mpi processes
- Take care that processes have enough resources, i.e., cores, assigned
- You want to make sure, that processes (and associated threads) do not move around, especially from one NUMA to another NUMA domain!

Example for starting with Open MPI :

```
$ mpiexec -n 2 --map-by slot:PE=2 --bind-to core -x OMP_NUM_THREADS=2 mpi+openmp_app ...
```

2 MPI processes


two slots, i.e., cores
per MPI process for
the threads

ensure MPI
processes do not
move around

pass environment variable
to MPI processes to tell
OpenMP to start 2 threads
per MPI process

Motivating MPI Partitioned Communication

Hybrid MPI+threads programs want to use fine grained communication to allow various optimizations




BUT: Individual send/recv operations for each thread come with overheads and prevent optimizations

MPI Partitioned P2P Communication

- Introduced in **MPI 4.0**
- **Extends** the **persistent P2P** communication
- Allows to **divide a message into a fixed number of *partitions*** for multiple contributions of data, potentially, by multiple actors (e.g. OpenMP threads)
- **Enables various optimizations:**
 - reduced MPI runtime overheads
 - message aggregation
 - relaxed ordering for partitions
- Adds **two new initialization methods** for send and receive, which cannot be intermixed with other send/recv methods


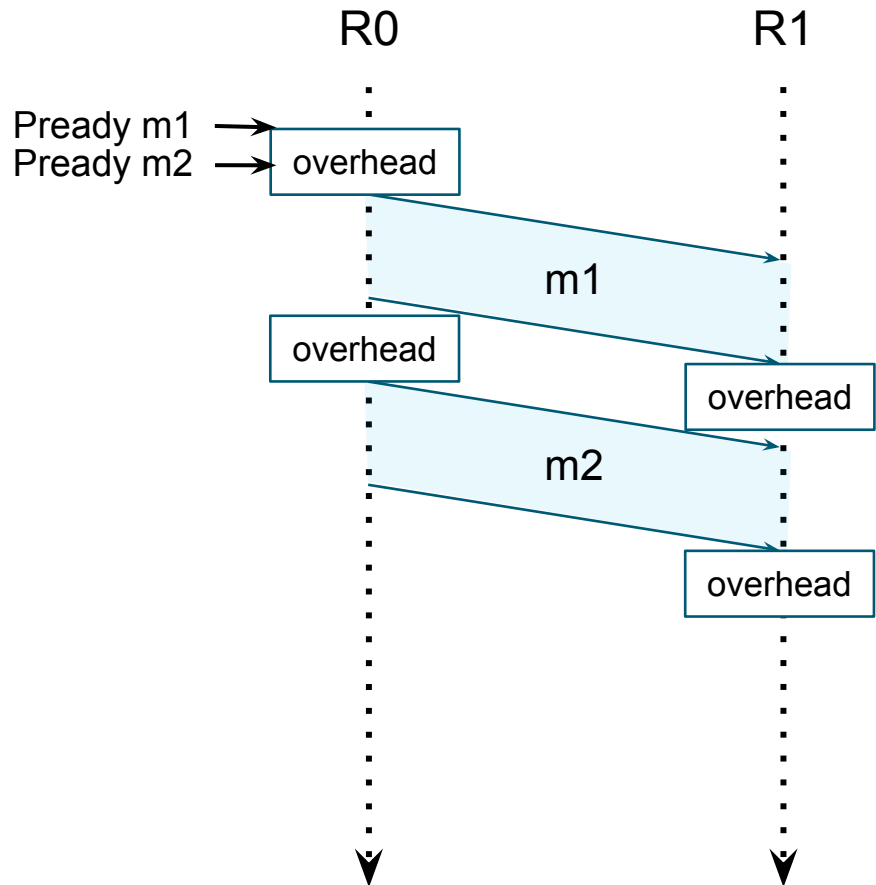
Possible optimizations: reduced runtime overheads

H L R I S




Possible optimizations: aggregation

H L R **I** S



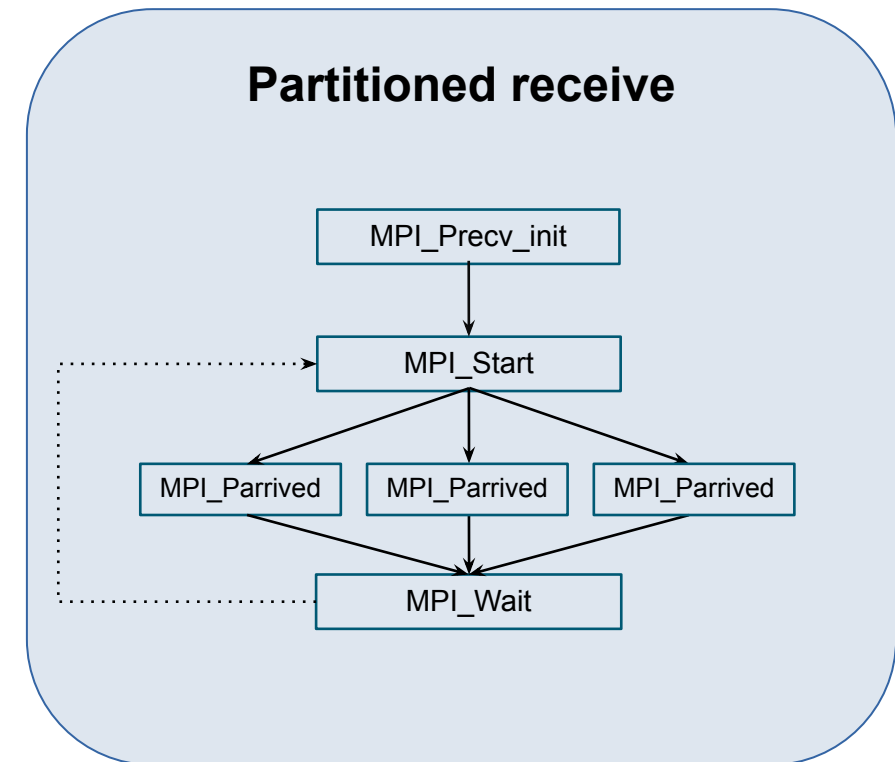
Partitioned communication API: sending



- **MPI_Psend_init:**
Initialize partitioned send operation for a buffer with a given number of partitions, returning a partitioned request handle.
- **MPI_Start:**
Start partitioned communication associated with a partitioned communication request.
- **MPI_Pready:**
Mark individual partition as ready to be sent. To be called for every partition.
- **MPI_Wait:**
Complete send operation (one round of sending).

Partitioned communication API: receiving

- **MPI_Psend_init:**
Initialize partitioned receive operation for a buffer with a given number of partitions, returning a partitioned request handle.
- **MPI_Start:**
Start partitioned communication associated with a partitioned communication request.
- **MPI_Parrived:**
Check if an individual partition is already received. To be called for every partition.
- **MPI_Wait:**
Complete receive operation

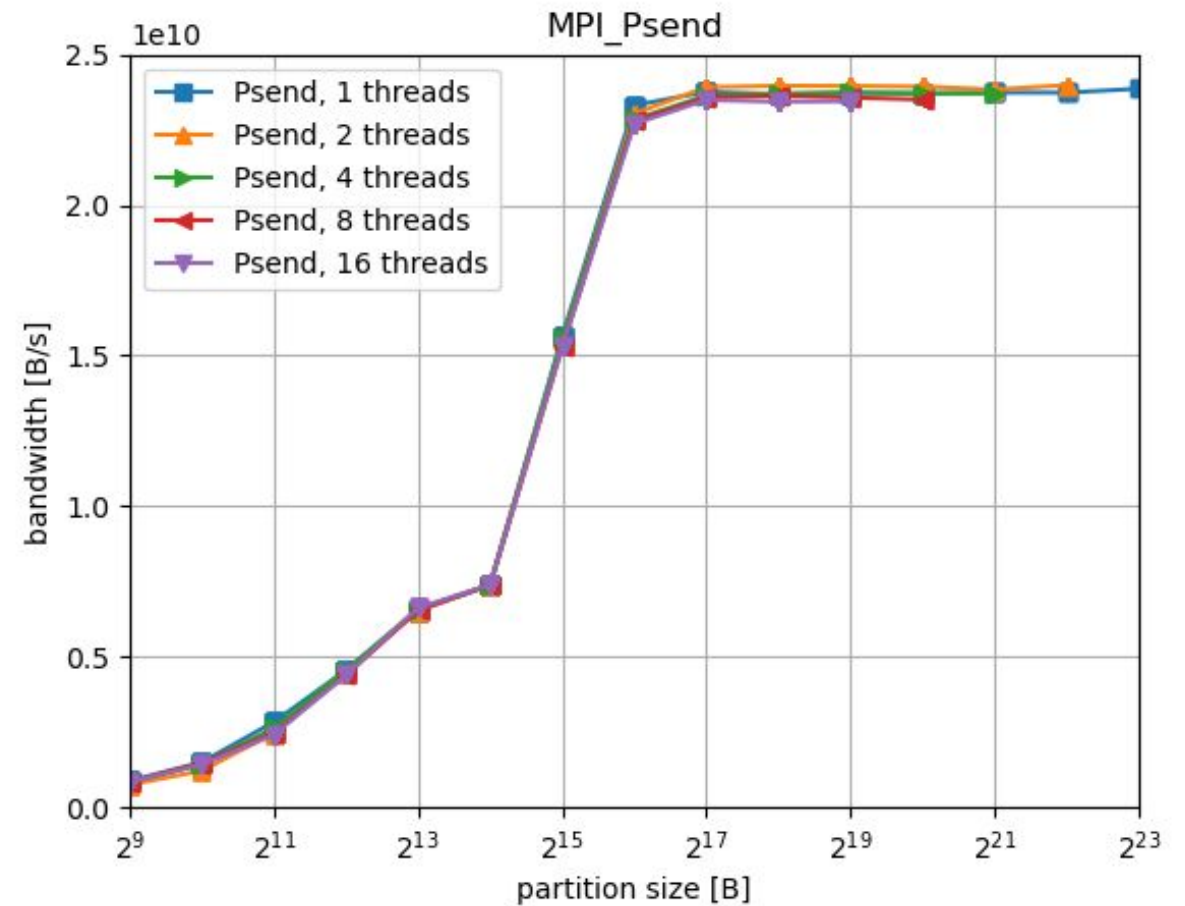



MPI Partitioned example with OpenMP

```
MPI_Psend_init( message,
                partitions_1, COUNT_1,
                MPI_DOUBLE, dest, tag,
                MPI_COMM_WORLD,
                MPI_INFO_NULL, &request );
MPI_Start(& request );
#pragma omp parallel for shared(request)
for ( i = 0; i < partitions ; ++ i ) {
    /* compute and fill partition #i */
    MPI_Pready(i , request );
}
while (! flag ) {
    MPI_Test( &request, &flag,
              MPI_STATUS_IGNORE );
    /* do useful work */
}
MPI_Request_free(& request );
```

```
MPI_Precv_init( message,
                partitions_2, COUNT_2,
                MPI_DOUBLE, src, tag,
                MPI_COMM_WORLD,
                MPI_INFO_NULL, &request );
MPI_Start(& request );
#pragma omp parallel for shared(request)
for ( i = 0; i < partitions ; ++ i ) {
    int part_flag = 0;
    while( 0 == part_flag ) {
        /* do something useful */
        MPI_Parrived(request, i, &part_flag );
    }
}
while (! flag ) {
    MPI_Test( &request, &flag,
              MPI_STATUS_IGNORE );
    /* do useful work */
}
MPI_Request_free(& request );
```

Multithreaded communication performance



A. Schneewind, C.Niethammer – EuroMPI 24 – Benchmarking the State of MPI Partitioned Communication in Open MPI

Message Passing Interface (MPI)

One-sided communication / RMA

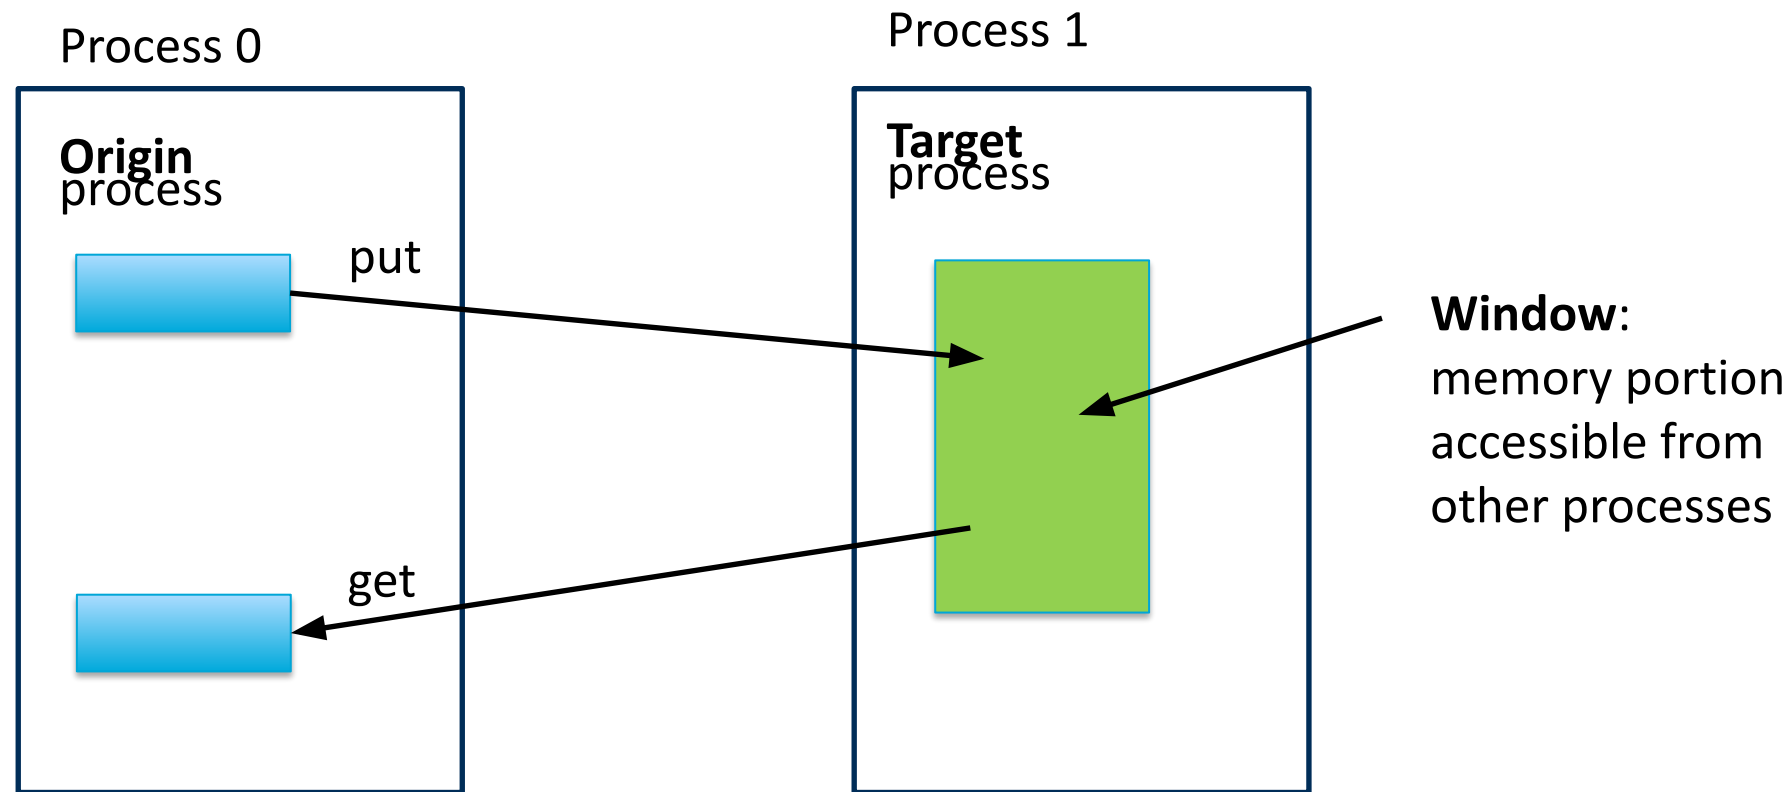
MPI one-sided communication

One-sided communication separates communication and synchronization

- **Reduces synchronization**
Put/Get operations focus only on transfer without synchronization
Only single process involved in transfer
Models directly RDMA in the hardware
- **Designed with communication computation overlap in mind:**
Put/Get are non-blocking
- **Mitigate some scalability problems of P2P**
Frequently changing/unknown communication partners

MPI one-sided model

H L R **I** S



MPI one-sided operations

- **Window creation and allocation**

Each process in a group of processes (Communicator) defines a chunk of his memory (Window), which can be afterwards accessed by all other processes in the group

- **Remote Memory Access (RMA)**

Access to remote windows:

- put, get, accumulate, ...

- **Synchronization**

RMA routines are non-blocking and must be surrounded by synchronization to guarantee that

- RMA is locally and remotely finished
- necessary cache operations are implicitly done

MPI window allocation

- Already allocated memory in the application:
`MPI_Win_create`
- Create buffer and make it available as window:
`MPI_Win_allocate`
`MPI_Win_allocate_shared`
- Buffer must be allocated later, e.g., size yet unknown
`MPI_Win_create_dynamic`

MPI one-sided data transfer

General Memory transfer:

- MPI_Get
- MPI_Put (race conditions!)

Atomic operations:

- MPI_Accumulate
- MPI_Get_accumulate
- MPI_Fetch_and_op
- MPI_Compare_and_swap



Get/fetch executed before the operation

R-Versions: MPI_Rget, MPI_Rput, ...

→ request-based, only with passive synchronization

MPI_Put / MPI_Get

- Non-blocking calls
- Origin specifies arguments for both, origin and target
- Equivalent to P2P transfer via send/recv operations
- target buffer is at address
 $\text{target_addr} = \text{win_base_target} + \text{target_disp_origin} * \text{disp_unit_target}$

```
MPI_Put(const void *origin_addr, int origin_count,      MPI_Datatype
origin_datatype, int target_rank,
        MPI_Aint target_disp, int target_count,
        MPI_Datatype target_datatype, MPI_Win win)
```

```
MPI_Get(void *origin_addr, int origin_count,
        MPI_Datatype origin_datatype, int target_rank,
        MPI_Aint target_disp, target_count,
        MPI_Datatype target_datatype, MPI_Win win)
```

MPI one-sided synchronization

- **Active synchronization:**
- Origin and target are involved
 - MPI_Win_fence
 - MPI_Win_post, MPI_Win_start, MPI_Win_complete, MPI_Win_wait
- **Passive synchronization:**
- Only origin process calls synchronization functions, target is passive
 - MPI_Win_lock, MPI_Win_unlock

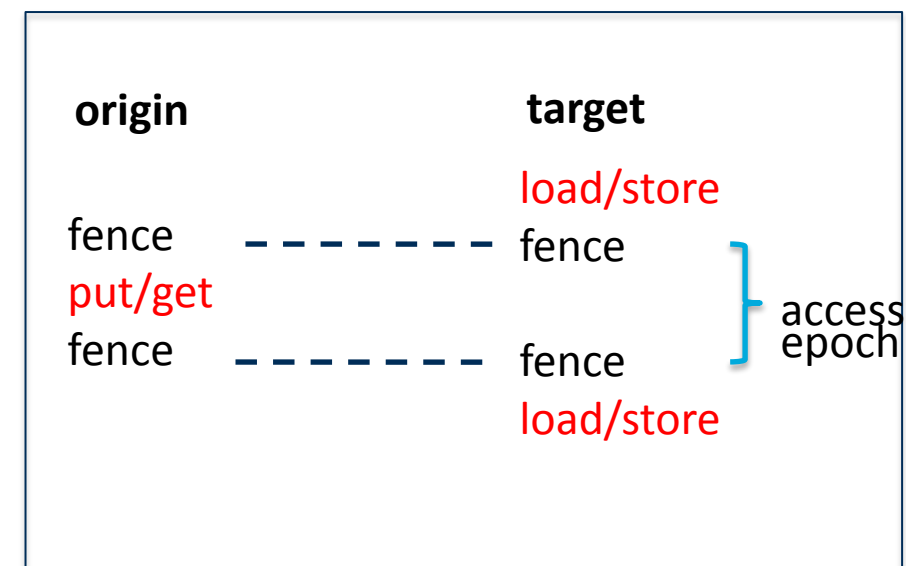
MPI active target synchronization

Barrier like synchronization is provided by

```
int MPI_Win_fence(int assert, MPI_Win win)
```

A fence ensures a consistent memory view before and after the fence for a provided window

Note: MPI_Barrier does not guarantee this!



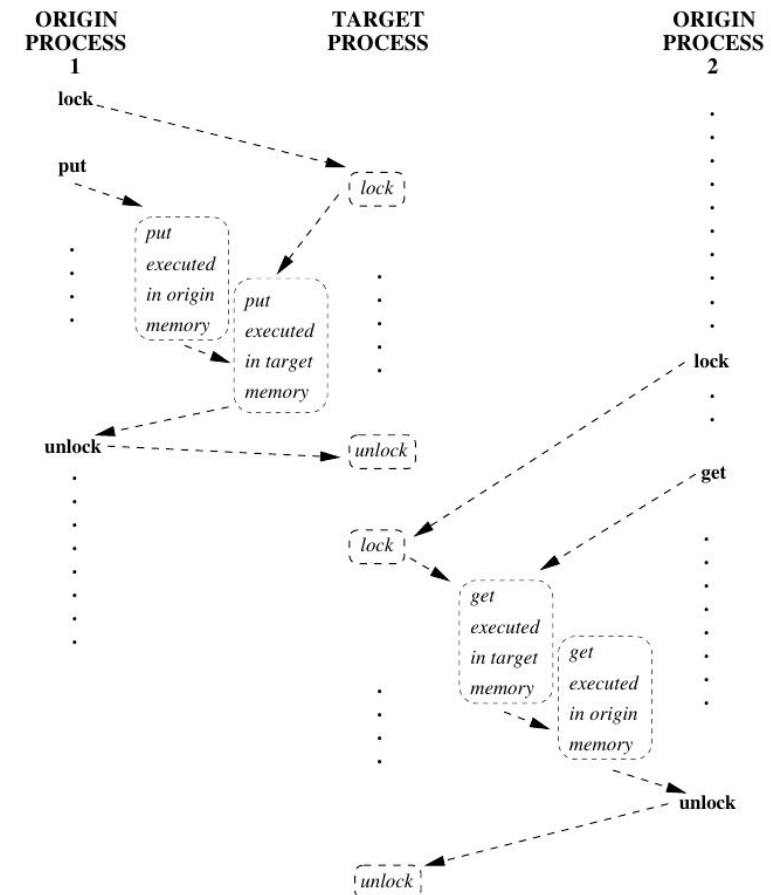
modify memory in window

MPI passive target synchronization

origin maintains locks for the target window via

```
int MPI_Win_lock(
    int lock_type,
    int rank, int assert,
    MPI_Win win)
```

```
int MPI_Win_unlock(
    int rank, MPI_Win win)
```



Literatur

- Gropp, B., Lusk, E.: Using MPI: Portable Parallel Programming with the Message-Passing Interface, 3rd Ed., 2014
- Ofizielle Versionen des Standards: <https://www.mpi-forum.org/docs/>