



# A Modular LLVM-Based Contract Framework for Parallel Programming Models

Yussur Mustafa Orajı

## Parallel Programming

- Working with language-native parallel programming difficult
  - Different language → different API
  - Low-level, little abstraction
- In contrast: Parallel Programming Models
  - MPI [7]
  - OpenSHMEM [3]
  - GASPI [6]
  - ...
- (Sometimes) support for multiple languages
- Very error-prone...

## How many issues can you spot in this tiny example?

```
#include <mpi.h>
#include <stdio.h>
int main (int argc, char** argv)
{
    int rank, size, buf[8];
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    MPI_Comm_size (MPI_COMM_WORLD, &size);

    MPI_Datatype type;
    MPI_Type_contiguous (2, MPI_INTEGER, &type);

    MPI_Recv (buf, 2, MPI_INT, size - rank, 123, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    MPI_Send (buf, 2, type, size - rank, 123, MPI_COMM_WORLD);
    printf ("Hello, I am rank %d of %d.\n", rank, size);

    return 0;
}
```

At least 8 issues in this code example!

## Our Approach

- Current example: MPI, but OpenSHMEM and other APIs are used as well
- Main languages: C/C++ and Fortran
- $\Rightarrow$  Need tool for each language and programming model. . .

## Our Approach

- Current example: MPI, but OpenSHMEM and other APIs are used as well
- Main languages: C/C++ and Fortran
- $\Rightarrow$  Need tool for each language and programming model. . .
- Idea: “Move the model out of the tool, apply analysis independent of model”

# Contract Basics

- Specify API assumptions in code
- Annotate additional information:
  - Preconditions / Premise: “What holds before this function?”
  - Postcondition: “What holds after it?”

<sup>1</sup>Feature delayed, but work continues [4]

## Contract Basics

- Specify API assumptions in code
- Annotate additional information:
  - Preconditions / Premise: “What holds before this function?”
  - Postcondition: “What holds after it?”
- Early C++20 draft<sup>1</sup>: `expects`, `ensures`, `later pre`, `post`
- Also available in other languages:
  - Prusti Project for Rust [1]
  - Kotlin Contracts [12]

---

<sup>1</sup>Feature delayed, but work continues [4]

## Applying Contracts

- Possible solutions using contracts (pseudocode):



## Applying Contracts

- Possible solutions using contracts (pseudocode):
  - `<ret-type> MPI_* PRE{ call!(MPI_Init) }  $\Rightarrow$  Ensures MPI_Init called`
  - `int MPI_Type_contiguous(..., MPI_Datatype* type) POST{  
no! (call!(MPI_Irecv,*type))  
until! (call!(MPI_Type_commit,type))}  $\Rightarrow$  Ensures type committed`

## Applying Contracts

- Possible solutions using contracts (pseudocode):
  - `<ret-type> MPI_* PRE{ call!(MPI_Init) }  $\Rightarrow$  Ensures MPI_Init called`
  - `int MPI_Type_contiguous(..., MPI_Datatype* type) POST{  
no! (call!(MPI_Irecv,*type))  
until! (call!(MPI_Type_commit,type))}  $\Rightarrow$  Ensures type committed`

**Analyses unaware of model: No need for additional analyses for other models!**

## Applying Contracts

- Possible solutions using contracts (pseudocode):
  - `<ret-type> MPI_* PRE{ call!(MPI_Init) }  $\Rightarrow$  Ensures MPI_Init called`
  - `int MPI_Type_contiguous(..., MPI_Datatype* type) POST{  
no! (call!(MPI_Irecv,*type))  
until! (call!(MPI_Type_commit,type))}  $\Rightarrow$  Ensures type committed`

**Analyses unaware of model: No need for additional analyses for other models!**

**Inherently extensible: For other error classes / other models, simply add the annotations!**

# Architecture Considerations

---


# Architecture Considerations

---

- For Proof of Concept:

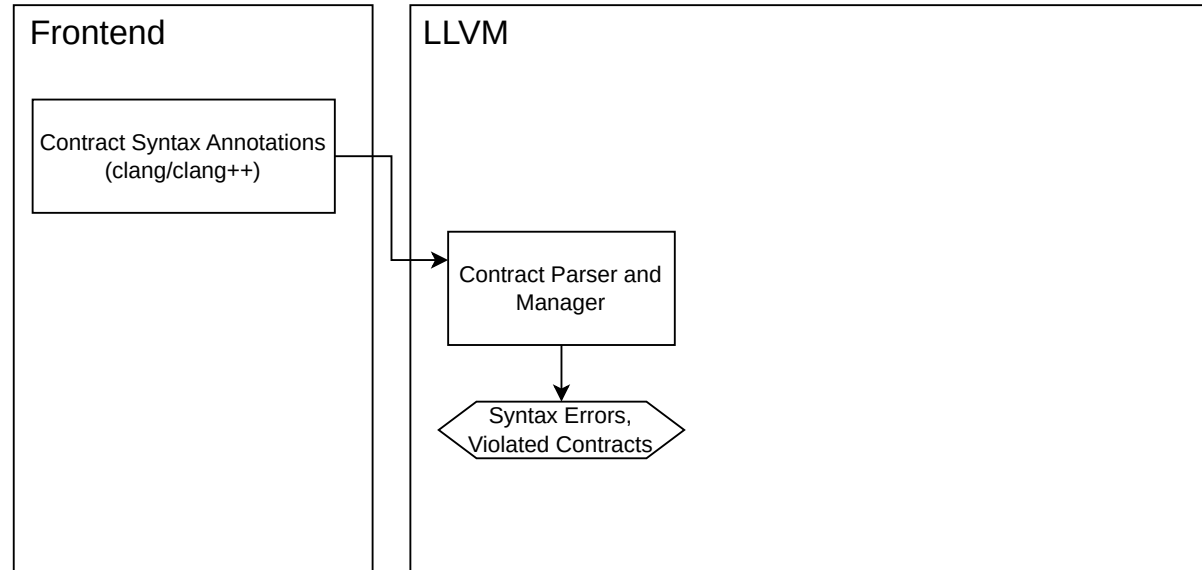
# Architecture Considerations

---



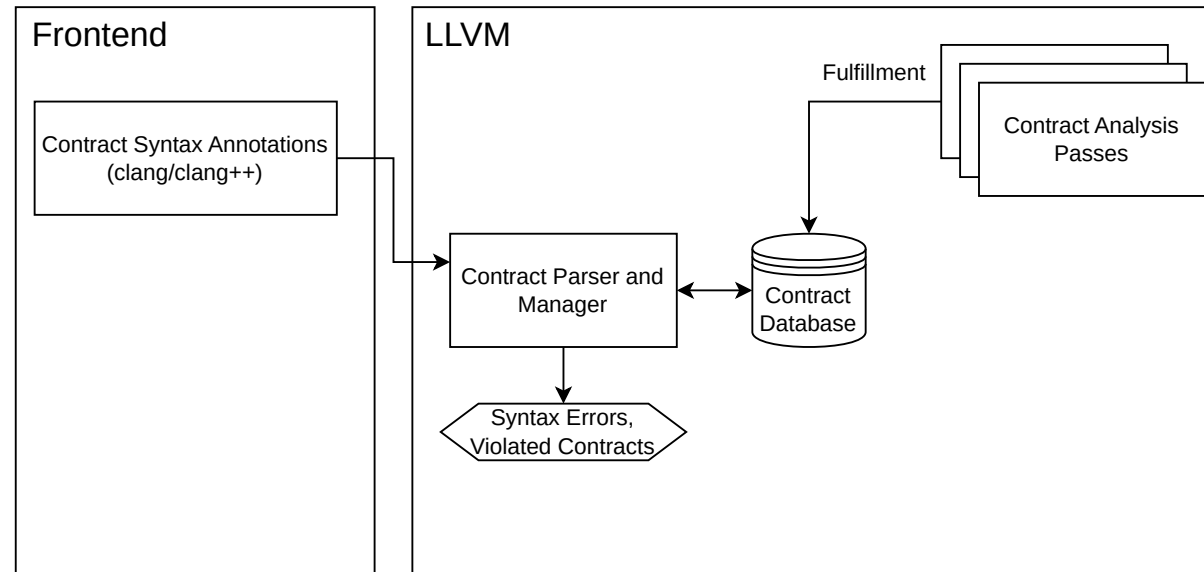
- For Proof of Concept:
  1. Develop contract syntax for considered languages  $\Rightarrow$  Function Annotations

# Architecture Considerations



- For Proof of Concept:
  1. Develop contract syntax for considered languages  $\Rightarrow$  Function Annotations
  2. Contract Parser and Manager as LLVM Analysis

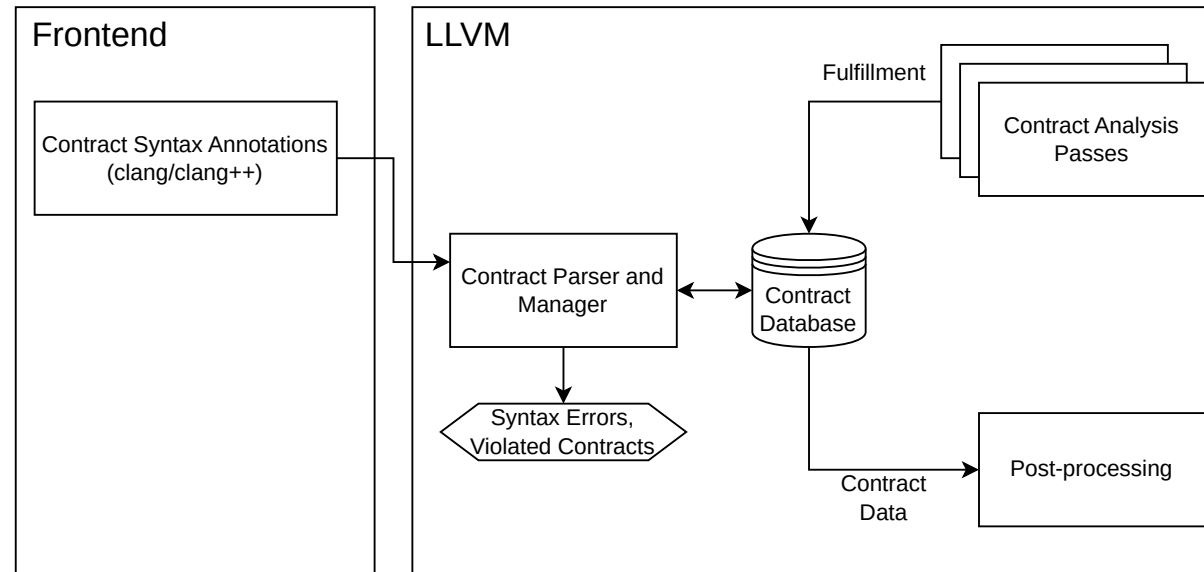
# Architecture Considerations



- For Proof of Concept:
  1. Develop contract syntax for considered languages  $\Rightarrow$  Function Annotations
  2. Contract Parser and Manager as LLVM Analysis
  3. Contract Analysis Pass(es) mark contracts as fulfilled, unable to prove, or violated in manager database

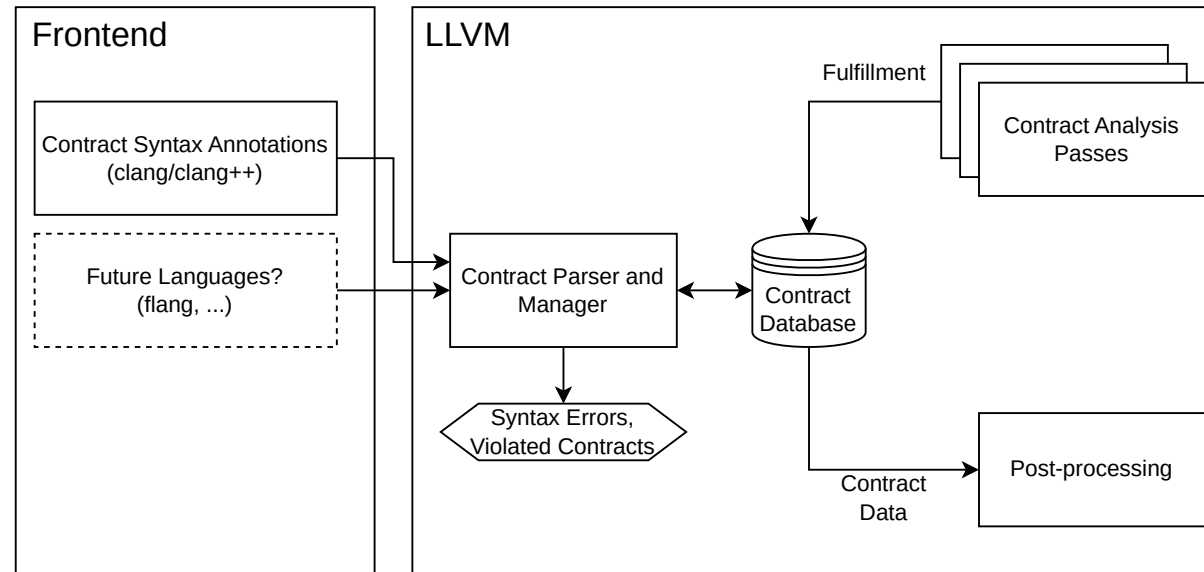


# Architecture Considerations



- For Proof of Concept:
  1. Develop contract syntax for considered languages  $\Rightarrow$  Function Annotations
  2. Contract Parser and Manager as LLVM Analysis
  3. Contract Analysis Pass(es) mark contracts as fulfilled, unable to prove, or violated in manager database
  4. Post-processing performs last steps before reporting issues if present

# Architecture Considerations



- For Proof of Concept:
  1. Develop contract syntax for considered languages  $\Rightarrow$  Function Annotations
  2. Contract Parser and Manager as LLVM Analysis
  3. Contract Analysis Pass(es) mark contracts as fulfilled, unable to prove, or violated in manager database
  4. Post-processing performs last steps before reporting issues if present

## Error Classes and Detection Methods

- Currently focusing on the following errors:

# Error Classes and Detection Methods

---

- Currently focusing on the following errors:
  1. Request Lifecycle errors: Request not checked or request leaked
    - Request reuse before completion leaks memory

```
1 MPI_Request req;  
2 MPI_Isend(&data, ..., &req);  
3 MPI_Isend(&data, ..., &req);  
4 // Only 2nd call is completed!  
5 MPI_Wait(&req, MPI_STATUS_IGNORE);
```

# Error Classes and Detection Methods

---

- Currently focusing on the following errors:
  1. Request Lifecycle errors: Request not checked or request leaked
  2. Resource (Datatype, Communicator) not initialized or freed
    - Many kinds of resources must be managed by the user manually
    - If not freed, data leaks occur. If not initialized, crashes or undefined behavior

```
1 MPI_Datatype type;  
2 MPI_Type_contiguous(..., &type);  
3 MPI_Send(&data, 1, type, ...);
```

# Error Classes and Detection Methods

---

- Currently focusing on the following errors:
  1. Request Lifecycle errors: Request not checked or request leaked
  2. Resource (Datatype, Communicator) not initialized or freed
  3. Local Data Race
    - Buffers of nonblocking operations may not be written to, and sometimes also not read

```
1 int value;  
2 MPI_Win_fence(0, win);  
3 if (rank == 0) {  
4     MPI_Get(&value, 1, ..., win);  
5     printf("Value is: %d", value);  
6 }  
7 MPI_Win_fence(0, win);
```

# Error Classes and Detection Methods

---

- Currently focusing on the following errors:
  1. Request Lifecycle errors: Request not checked or request leaked
  2. Resource (Datatype, Communicator) not initialized or freed
  3. Local Data Race
  4. Missing Init/Finalize

# Error Classes and Detection Methods

---

- Currently focusing on the following errors:
  1. Request Lifecycle errors: Request not checked or request leaked
  2. Resource (Datatype, Communicator) not initialized or freed
  3. Local Data Race
  4. Missing Init/Finalize
  5. Unmatched wait, Mixed RMA sync or outside epoch, ...



# Error Classes and Detection Methods

---

- Currently focusing on the following errors:
  1. Request Lifecycle errors: Request not checked or request leaked
  2. Resource (Datatype, Communicator) not initialized or freed
  3. Local Data Race
  4. Missing Init/Finalize
  5. Unmatched wait, Mixed RMA sync or outside epoch, ...
- Error classification leads to planned contract syntax
  - Require function called before/after current: `PRE / POST { call!(f) }`
  - Require a release operator: `POST { no! (op) until! (call!(f)) }`

# Error Classes and Detection Methods

---

- Currently focusing on the following errors:
  1. **Request Lifecycle errors: Request not checked or request leaked:** `POST { call!(MPI_Wait) }`
  2. Resource (Datatype, Communicator) not initialized or freed
  3. Local Data Race
  4. **Missing Init/Finalize:** `PRE { call!(MPI_Init) }`
  5. Unmatched wait, Mixed RMA sync or outside epoch, ...
  
- Error classification leads to planned contract syntax
  - **Require function called before/after current:** `PRE / POST { call!(f) }`
  - Require a release operator: `POST { no! (op) until! (call!(f)) }`

# Error Classes and Detection Methods

---

- Currently focusing on the following errors:
  1. Request Lifecycle errors: Request not checked or request leaked
  2. Resource (Datatype, Communicator) not initialized or freed
  3. **Local Data Race**: `POST { no! (read!(buf)) until! (call!(MPI_Win_fence)) }`
  4. Missing Init/Finalize
  5. Unmatched wait, Mixed RMA sync or outside epoch, ...
- Error classification leads to planned contract syntax
  - Require function called before/after current: `PRE / POST { call!(f) }`
  - **Require a release operator**: `POST { no! (op) until! (call!(f)) }`

## Contract Abstraction

- Use ANTLR [14] to define formal grammar
- C/C++: Define contracts as function annotations
  - Helper macro: `#define CONTRACT(x) __attribute__((annotate("CONTRACT{" #x "}")))`

## Contract Abstraction

- Use ANTLR [14] to define formal grammar
- C/C++: Define contracts as function annotations
  - Helper macro: `#define CONTRACT(x) __attribute__((annotate("CONTRACT{" #x "}")))`

Annotating using the helper macro:

```
1 int MPI_Finalize(void)
2 CONTRACT(PRE {call!(MPI_Init)});
```

## Contract Abstraction

- Use ANTLR [14] to define formal grammar
- C/C++: Define contracts as function annotations
  - Helper macro: `#define CONTRACT(x) __attribute__((annotate("CONTRACT{" #x "}")))`
- Allow defining tags when functions fulfill similar purpose: `MPI_Win_fence`  $\rightarrow$  TAGS { `rma_complete(1)` }

Annotating using the helper macro:

```
1 int MPI_Finalize(void)
2 CONTRACT(PRE {call_tag!(initialize)});

1 int MPI_Init(...) CONTRACT (TAGS { initialize })
2 int MPI_Init_thread(...) CONTRACT (TAGS { initialize })
```

## Contract Abstraction

- Use ANTLR [14] to define formal grammar
- C/C++: Define contracts as function annotations
  - Helper macro: `#define CONTRACT(x) __attribute__((annotate("CONTRACT{" #x "}")))`
- Allow defining tags when functions fulfill similar purpose: `MPI_Win_fence`  $\rightarrow$  TAGS { `rma_complete(1)` }

Annotating using the helper macro:

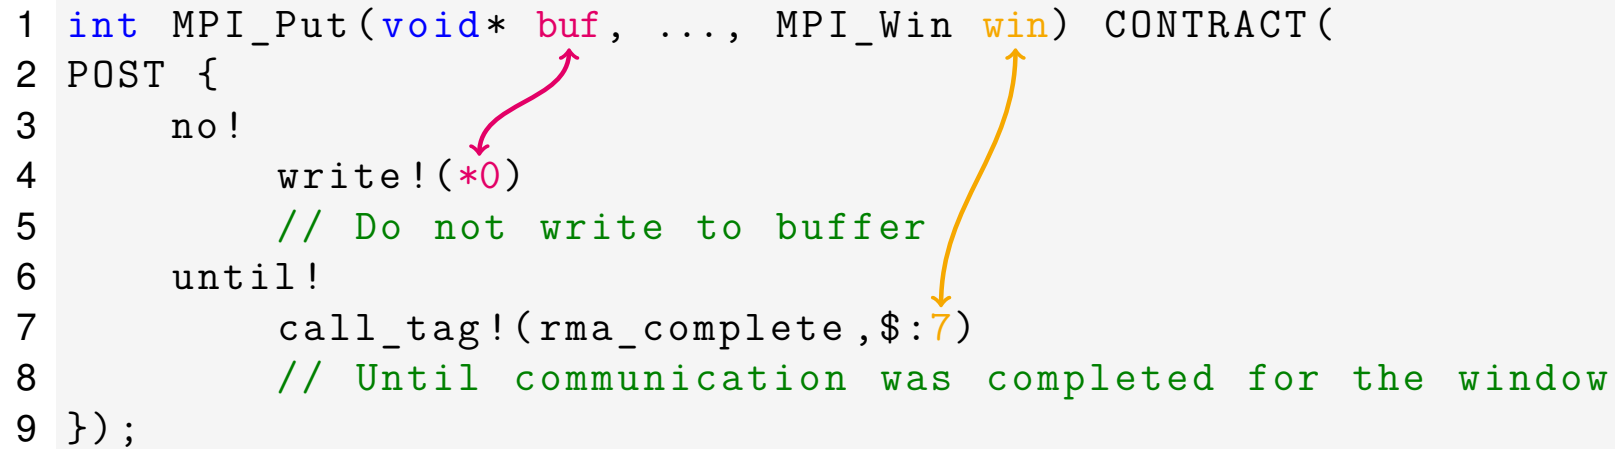
```
1 int MPI_Put(void* buf, ..., MPI_Win win) CONTRACT(  
2 POST {  
3     no!  
4         write!(*0)  
5         // Do not write to buffer  
6     until!  
7         call_tag!(rma_complete,$:7)  
8         // Until communication was completed for the window  
9 });
```

## Contract Abstraction

- Use ANTLR [14] to define formal grammar
- C/C++: Define contracts as function annotations
  - Helper macro: `#define CONTRACT(x) __attribute__((annotate("CONTRACT{" #x "}")))`
- Allow defining tags when functions fulfill similar purpose: `MPI_Win_fence`  $\rightarrow$  TAGS { `rma_complete(1)` }

Annotating using the helper macro:

```
1 int MPI_Put(void* buf, ..., MPI_Win win) CONTRACT(  
2 POST {  
3     no!  
4     write!(*0)  
5     // Do not write to buffer  
6     until!  
7     call_tag!(rma_complete, $:7)  
8     // Until communication was completed for the window  
9 });
```





## Contract Abstraction

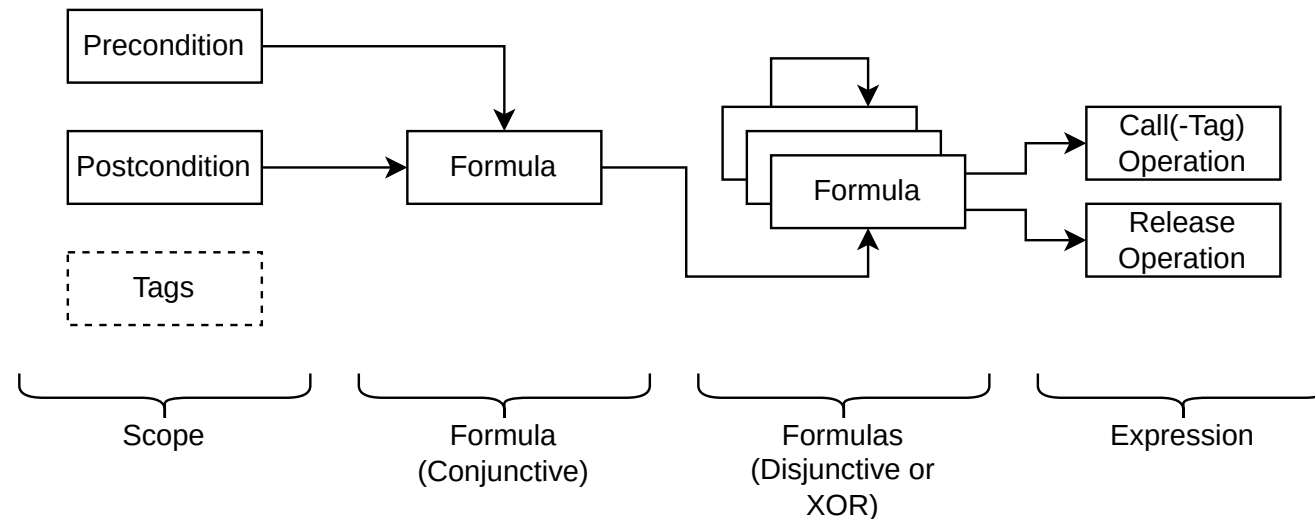
- Use ANTLR [14] to define formal grammar
- C/C++: Define contracts as function annotations
  - Helper macro: `#define CONTRACT(x) __attribute__((annotate("CONTRACT{" #x "}")))`
- Allow defining tags when functions fulfill similar purpose: `MPI_Win_fence`  $\rightarrow$  TAGS { `rma_complete(1)` }
- All contracts may be written at *any* function declaration:
  - Modify original model header (e.g. `mpi.h`)
  - Add additional declarations as separate header
  - Add declarations in source for specific analyses only
  - Chef's Choice: Separate header, include using `-include` preprocessor flag

## Contract Abstraction

- Contract Manager: LLVM Analysis that reads `@llvm.global.annotations` → List of annotations
  - Calls parser, transform contract string into `ContractTree` representation
  - $\Rightarrow$  Analyses only interact with `ContractTree`, no contract language details needed

## Contract Abstraction

- Contract Manager: LLVM Analysis that reads `@llvm.global.annotations` → List of annotations
  - Calls parser, transform contract string into `ContractTree` representation
  - $\Rightarrow$  Analyses only interact with `ContractTree`, no contract language details needed



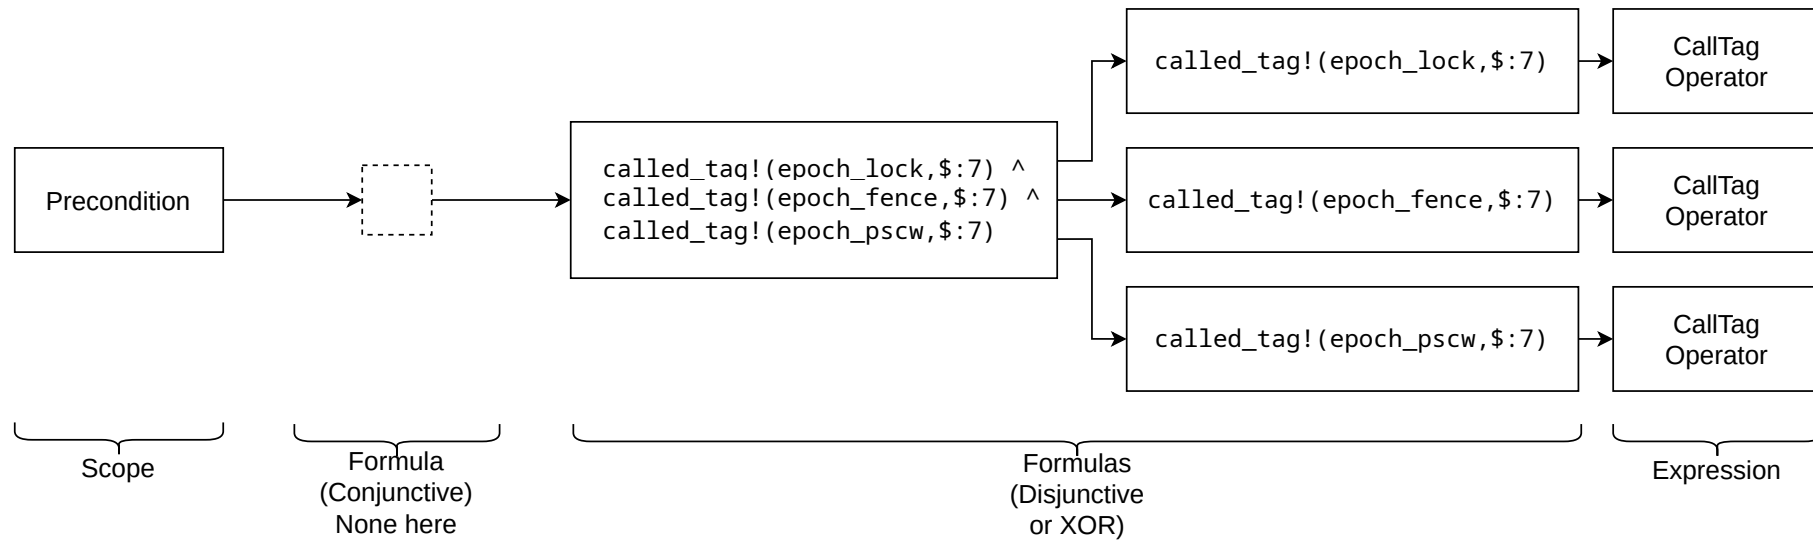
## Contract Abstraction

- Contract Manager: LLVM Analysis that reads `@llvm.global.annotations` → List of annotations
  - Calls parser, transform contract string into `ContractTree` representation
  - $\Rightarrow$  Analyses only interact with `ContractTree`, no contract language details needed
- Example: “For an RMA operation, an epoch should exist (any kind, but not mixed)”

```
1 int MPI_Put(...) CONTRACT (  
2 PRE {  
3     call_tag!(epoch_lock,$:7)} ^  
4     call_tag!(epoch_fence,$:7)^  
5     call_tag!(epoch_pscw,$:7)  
6 });
```

## Contract Abstraction

- Contract Manager: LLVM Analysis that reads `@llvm.global.annotations` → List of annotations
  - Calls parser, transform contract string into `ContractTree` representation
  - $\Rightarrow$  Analyses only interact with `ContractTree`, no contract language details needed
- Example: “For an RMA operation, an epoch should exist (any kind, but not mixed)”



## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: Check if function called previously
  - PostCall: Check if function called afterward
  - Release: Check for conflict before release
- Verifiers share generic worklist algorithm with interprocedural support

## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: **Check if function called previously**: `MPI_Put` → `PRE { call!(MPI_Win_fence) }`
  - PostCall: Check if function called afterward
  - Release: Check for conflict before release
- Verifiers share generic worklist algorithm with interprocedural support. Current: `{(1, 2), (1, 4)}`

```
1 if (...)
2     MPI_Win_fence(win)
3     x++
4 print(x)
5 MPI_Put(&x, ..., win)
6 x = 1
7 return 0
```

Line	Result
1	<b>NOTCALLED</b>
2	CALLED
3	CALLED
4	CALLED
5	CALLED
6	CALLED
7	CALLED

## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: **Check if function called previously**: `MPI_Put` → `PRE { call!(MPI_Win_fence) }`
  - PostCall: Check if function called afterward
  - Release: Check for conflict before release
- Verifiers share generic worklist algorithm with interprocedural support. Current: `{(2, 3), (1, 4)}`

```
1  if (...)
2      MPI_Win_fence(win)
3      x++
4  print(x)
5  MPI_Put(&x, ..., win)
6  x = 1
7  return 0
```

Line	Result
1	NOTCALLED
2	<b>PARAMCHECK</b>
3	CALLED
4	CALLED
5	CALLED
6	CALLED
7	CALLED



## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: **Check if function called previously**: `MPI_Put`  $\rightarrow$  `PRE { call!(MPI_Win_fence) }`
  - PostCall: Check if function called afterward
  - Release: Check for conflict before release
- Verifiers share generic worklist algorithm with interprocedural support. Current: `{(4, 5), (1, 4)}`

```
1 if (...)
2     MPI_Win_fence(win)
3     x++
4 print(x)
5 MPI_Put(&x, ..., win)
6 x = 1
7 return 0
```

Line	Result
1	NOTCALLED
2	PARAMCHECK
3	<b>PARAMCHECK</b>
4	<b>PARAMCHECK</b>
5	CALLED
6	CALLED
7	CALLED

## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: **Check if function called previously**: `MPI_Put` → `PRE { call!(MPI_Win_fence) }`
  - PostCall: Check if function called afterward
  - Release: Check for conflict before release
- Verifiers share generic worklist algorithm with interprocedural support. Current: `{(4, 5)}`

```
1 if (...)
2     MPI_Win_fence(win)
3     x++
4 print(x)
5 MPI_Put(&x, ..., win)
6 x = 1
7 return 0
```

Line	Result
1	NOTCALLED
2	PARAMCHECK
3	PARAMCHECK
4	<b>NOTCALLED</b>
5	CALLED
6	CALLED
7	CALLED

## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: **Check if function called previously**: `MPI_Put` → `PRE { call!(MPI_Win_fence) }`
  - PostCall: Check if function called afterward
  - Release: Check for conflict before release
- Verifiers share generic worklist algorithm with interprocedural support. Current: `∅`

```
1 if (...)
2     MPI_Win_fence(win)
3     x++
4 print(x)
5 MPI_Put(&x, ..., win)
6 x = 1
7 return 0
```

Line	Result
1	NOTCALLED
2	PARAMCHECK
3	PARAMCHECK
4	NOTCALLED
5	<b>ERROR</b>
6	<b>ERROR</b>
7	<b>ERROR</b>

## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: **Check if function called previously**: `MPI_Put` → `PRE { call!(MPI_Win_fence) }`
  - PostCall: Check if function called afterward
  - Release: Check for conflict before release
- Verifiers share generic worklist algorithm with interprocedural support

```
1 if (...)
2     MPI_Win_fence(win)
3     x++
4 print(x)
5 MPI_Put(&x, ..., win)
6 x = 1
7 return 0
```

Line	Result
1	NOTCALLED
2	PARAMCHECK
3	PARAMCHECK
4	NOTCALLED
5	ERROR
6	ERROR
7	ERROR

## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: Check if function called previously
  - PostCall: Check if function called afterward
  - Release: **Check for conflict before release**: `MPI_Put → POST { no! (write!(*0)) until! ...}`
- Verifiers share generic worklist algorithm with interprocedural support . Current: `{(5, 6)}`

```
1 if (...)
2     MPI_Win_fence(win)
3     x++
4 print(x)
5 MPI_Put(&x, ..., win)
6 x = 1
7 return 0
```

Line	Result
1	FULFILLED
2	FULFILLED
3	FULFILLED
4	FULFILLED
5	<b>FORBIDDEN</b>
6	FULFILLED
7	FULFILLED

## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: Check if function called previously
  - PostCall: Check if function called afterward
  - Release: **Check for conflict before release**: `MPI_Put → POST { no! (write!(*0)) until! ...}`
- Verifiers share generic worklist algorithm with interprocedural support . Current:  $\emptyset$

```
1  if (...)
2      MPI_Win_fence(win)
3      x++
4  print(x)
5  MPI_Put(&x, ..., win)
6  x = 1
7  return 0
```

Line	Result
1	FULFILLED
2	FULFILLED
3	FULFILLED
4	FULFILLED
5	FORBIDDEN
6	<b>ERROR</b>
7	<b>ERROR</b>

## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: Check if function called previously
  - PostCall: Check if function called afterward
  - Release: **Check for conflict before release**: `MPI_Put → POST { no! (write!(*0)) until! ...}`
- Verifiers share generic worklist algorithm with interprocedural support

```
1 if (...)
2     MPI_Win_fence(win)
3     x++
4 print(x)
5 MPI_Put(&x, ..., win)
6 x = 1
7 return 0
```

Line	Result
1	FULFILLED
2	FULFILLED
3	FULFILLED
4	FULFILLED
5	FORBIDDEN
6	ERROR
7	ERROR

## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: Check if function called previously
  - PostCall: Check if function called afterward
  - Release: Check for conflict before release
- Verifiers share generic worklist algorithm with interprocedural support
- Post-processing pass outputs results for each contract
  - First: Resolve Formulas
  - Then: Print state, debug info of contract if violated or unknown



## Analysis

- Contract Verifiers: Implemented using Data Flow Analysis
  - PreCall: Check if function called previously
  - PostCall: Check if function called afterward
  - Release: Check for conflict before release
- Verifiers share generic worklist algorithm with interprocedural support

```
1  ## Contract violation detected! ##
2  --> Function: MPI_Put
3  --> Contract: CONTRACT{POST { no! (write!(*0)) ...}}
4  --> Postcondition Subformula Status: Violated
5      --> Formula String: no!(write!(*0))...
6      --> Error Info:
7          [ContractVerifierRelease] Found store at main:6 which
              is in conflict with MPI_Put at main:5 before release
```

## Experiment Setup

- Evaluate both accuracy and compilation overhead
- Accuracy:
  - Tests from MPI-BugBench (Level 1) [10]
  - Tests from RMA RaceBench [16]
- Compilation Overhead:
  - LULESH [11]
  - miniVite [8]
  - PRK Stencil [17]
- Compare against PARCOACH-static [15] and MPI-Checker [5]

## Accuracy

	Local Data Races (RaceBench)	Accuracy
CoVer	<div><div>20</div><div>18</div><div>8</div></div>	0.809
PARCOACH	<div><div>19</div><div>10</div><div>9</div><div>9</div></div>	0.617
MPI-Checker	<div><div>19</div><div>28</div></div>	0.404




	Local Data Races (BugBench)	Accuracy
CoVer	<div><div>8</div><div>7</div><div>1</div></div>	0.938
PARCOACH	<div><div>4</div><div>3</div><div>5</div><div>4</div></div>	0.438
MPI-Checker	<div><div>7</div><div>2</div><div>7</div></div>	0.438



- Focus on only Local Races (direct comparison to PARCOACH)
- Significantly reduced FP rate

## Accuracy

	Common RMA Issues	Accuracy
CoVer	<div><div>5</div><div>1</div></div>	0.833
PARCOACH	<div>6</div>	0.000
MPI-Checker	<div>6</div>	0.000



	Request Lifecycle	Accuracy
CoVer	<div>16</div>	1.000
PARCOACH	<div>16</div>	0.000
MPI-Checker	<div>15</div> <div>1</div>	0.938



- Focus on only Local Races (direct comparison to PARCOACH)
- Significantly reduced FP rate
- CoVer: Wide error class coverage
  - Adding checks often trivial
- CoVer and MPI-Checker detect similar request lifecycle issues

## Accuracy

	RMARaceBench	Accuracy
CoVer	<div><div>20</div><div>43</div><div>19</div><div>43</div></div>	0.504
PARCOACH	<div><div>19</div><div>34</div><div>34</div><div>38</div></div>	0.424
MPI-Checker	<div><div>53</div><div>72</div></div>	0.424




	MPI-BugBench	Accuracy
CoVer	<div><div>30</div><div>22</div><div>75</div></div>	0.400
PARCOACH	<div><div>4</div><div>14</div><div>25</div><div>87</div></div>	0.138
MPI-Checker	<div><div>15</div><div>22</div><div>91</div></div>	0.280



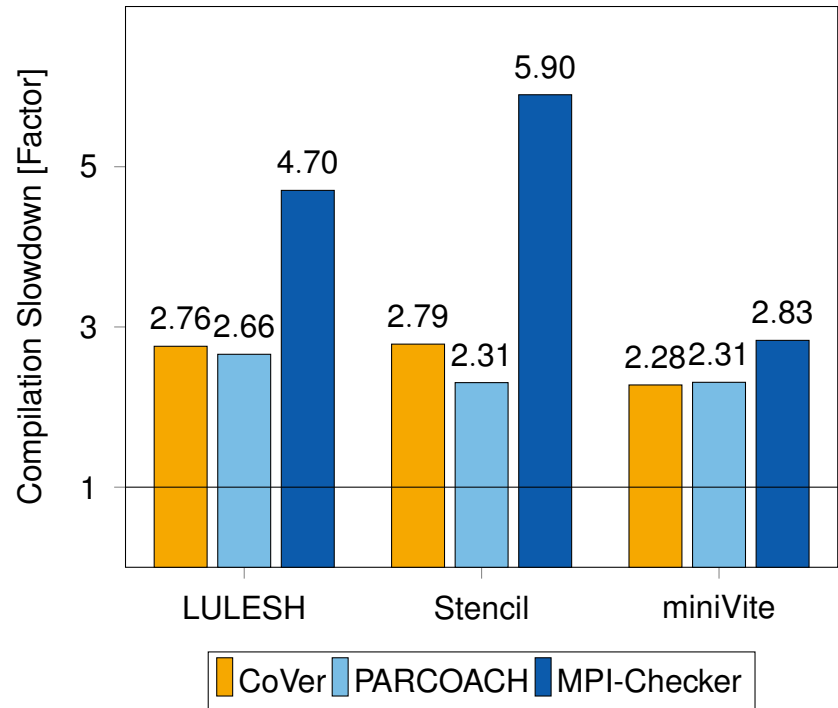
- Focus on only Local Races (direct comparison to PARCOACH)
- Significantly reduced FP rate
- CoVer: Wide error class coverage
  - Adding checks often trivial
- CoVer and MPI-Checker detect similar request lifecycle issues
- Overall: High Accuracy using wide coverage
- But: (In total) Higher FP rate than MPI-Checker
  - Caused by ignoring branch condition

## Overhead




- PARCOACH and CoVer have similar runtime

## Overhead



- PARCOACH and CoVer have similar runtime
- Note: *All* tools print false positives for these applications
  - Causes: Unknown branch conditions, function pointers, ...


## Overhead



- PARCOACH and CoVer have similar runtime
- Note: *All* tools print false positives for these applications
  - Causes: Unknown branch conditions, function pointers, ...
- MPI-Checker prints the least → Advantage of symbolic execution



## Overhead



- PARCOACH and CoVer have similar runtime
- Note: *All* tools print false positives for these applications
  - Causes: Unknown branch conditions, function pointers, ...
- MPI-Checker prints the least → Advantage of symbolic execution
  - Disadvantage shown in the graph...

## Related Work

---

- MUST [9], PARCOACH [15], ITAC [13] are dynamic MPI error detection tools
  - Significant runtime overhead
  - PARCOACH can run statically, but has no support for interprocedural contexts

## Related Work

---

- MUST [9], PARCOACH [15], ITAC [13] are dynamic MPI error detection tools
  - Significant runtime overhead
  - PARCOACH can run statically, but has no support for interprocedural contexts
- MPI-Checker [5] is a static correctness checker
  - Symbolic Execution
  - No support across translation units
  - Bound to `clang` frontend

## Related Work

---

- MUST [9], PARCOACH [15], ITAC [13] are dynamic MPI error detection tools
  - Significant runtime overhead
  - PARCOACH can run statically, but has no support for interprocedural contexts
- MPI-Checker [5] is a static correctness checker
  - Symbolic Execution
  - No support across translation units
  - Bound to `clang` frontend
- SPMD IR [2] is an approach unifying multiple programming models
  - Generic, but non-modular: Cannot extend without source modification
  - Both error classes and supported models static once tool shipped

# Conclusion

---

## Tool Summary:

- Ensure modularity by removing model info from analysis
- Use contracts to verify program properties
- Can detect a wide variety of issues, easy to extend

# Conclusion

## Tool Summary:

- Ensure modularity by removing model info from analysis
- Use contracts to verify program properties
- Can detect a wide variety of issues, easy to extend

How many issues can you spot in this tiny example?

```
#include <mpi.h>
#include <stdio.h>
int main (int argc, char** argv)
{
    int rank, size, buf[8];
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    MPI_Comm_size (MPI_COMM_WORLD, &size);

    MPI_Datatype type;
    MPI_Type_contiguous (2, MPI_INTEGER, &type);

    MPI_Recv (buf, 2, MPI_INT, size - rank, 123, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    MPI_Send (buf, 2, type, size - rank, 123, MPI_COMM_WORLD);
    printf ("Hello, I am rank %d of %d.\n", rank, size);
    return 0;
}
```

At least 8 issues in this code example!

MPI + GPU Correctness Checking with MUST  
46th VI-HPS Workshop  
Joachim Jenke, Felix Tomsik

NHR4  
CES  
NHR for  
Computational  
Engineering  
Science

pop  
Performance Optimisation  
and Productivity  
A Centre of Excellence in HPC

it2  
High  
Performance  
Computing

RWTH AACHEN  
UNIVERSITY

# Conclusion

## Tool Summary:

- Ensure modularity by removing model info from analysis
- Use contracts to verify program properties
- Can detect a wide variety of issues, easy to extend

## Evaluation Results:

- Wide coverage grants high accuracy
- Data Flow analysis comparatively cheap

How many issues can you spot in this tiny example?

```
#include <mpi.h>
#include <stdio.h>
int main (int argc, char** argv)
{
    int rank, size, buf[8];
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    MPI_Comm_size (MPI_COMM_WORLD, &size);

    MPI_Datatype type;
    MPI_Type_contiguous (2, MPI_INTEGER, &type);

    MPI_Recv (buf, 2, MPI_INT, size - rank, 123, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    MPI_Send (buf, 2, type, size - rank, 123, MPI_COMM_WORLD);
    printf ("Hello, I am rank %d of %d.\n", rank, size);
    return 0;
}
```

At least 8 issues in this code example!

MPI + GPU Correctness Checking with MUST  
46th VI-HPS Workshop  
Joachim Jenke, Felix Tomsik

NHR4  
CES  
NHR for  
Computational  
Engineering  
Science

pop  
Performance Optimisation  
and Productivity  
A Centre of Excellence in HPC

it2  
High  
Performance  
Computing

RWTH AACHEN  
UNIVERSITY

# Conclusion

## Tool Summary:

- Ensure modularity by removing model info from analysis
- Use contracts to verify program properties
- Can detect a wide variety of issues, easy to extend

## Evaluation Results:

- Wide coverage grants high accuracy
- Data Flow analysis comparatively cheap

## Future Work:

- Improve analysis accuracy
- Extend programming language support
- Extend contract language for further error classes

How many issues can you spot in this tiny example?

```
#include <mpi.h>
#include <stdio.h>
int main (int argc, char** argv)
{
    int rank, size, buf[8];
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    MPI_Comm_size (MPI_COMM_WORLD, &size);

    MPI_Datatype type;
    MPI_Type_contiguous (2, MPI_INTEGER, &type);

    MPI_Recv (buf, 2, MPI_INT, size - rank, 123, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    MPI_Send (buf, 2, type, size - rank, 123, MPI_COMM_WORLD);
    printf ("Hello, I am rank %d of %d.\n", rank, size);
    return 0;
}
```

At least 8 issues in this code example!

MPI + GPU Correctness Checking with MUST  
46th VI-HPS Workshop  
Joachim Jenke, Felix Tomsik

NHR4  
CES  
NHR for  
Computational  
Engineering  
Science

POP  
Performance Optimisation  
and Productivity  
A Centre of Excellence in HPC

it2  
High  
Performance  
Computing

RWTH AACHEN  
UNIVERSITY



# Conclusion

## Tool Summary:

- Ensure modularity by removing model info from analysis
- Use contracts to verify program properties
- Can detect a wide variety of issues, easy to extend

## Evaluation Results:

- Wide coverage grants high accuracy
- Data Flow analysis comparatively cheap

## Future Work:

- Improve analysis accuracy
- Extend programming language support
- Extend contract language for further error classes

How many issues can you spot in this tiny example?

```
#include <mpi.h>
#include <stdio.h>
int main (int argc, char** argv)
{
    int rank, size, buf[8];
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    MPI_Comm_size (MPI_COMM_WORLD, &size);

    MPI_Datatype type;
    MPI_Type_contiguous (2, MPI_INTEGER, &type);

    MPI_Recv (buf, 2, MPI_INT, size - rank, 123, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    MPI_Send (buf, 2, type, size - rank, 123, MPI_COMM_WORLD);
    printf ("Hello, I am rank %d of %d.\n", rank, size);
    return 0;
}
```

At least 8 issues in this code example!

MPI + GPU Correctness Checking with MUST  
46th VI-HPS Workshop  
Joachim Jenke, Felix Tomsik

NHR4  
CES  
NHR for  
Computational  
Engineering  
Science

POP  
Performance Optimisation  
and Productivity  
A Centre of Excellence in HPC

it2  
High  
Performance  
Computing

RWTH AACHEN  
UNIVERSITY

# Thank you for your attention!

## References I

---

- [1] Vytautas Astrauskas et al. “The Prusti Project: Formal Verification for Rust”. In: *NASA Formal Methods*. Cham: Springer International Publishing, 2022, pp. 88–108. ISBN: 978-3-031-06773-0. DOI: 10.1007/978-3-031-06773-0\_5.
- [2] Semih Burak et al. “SPMD IR: Unifying SPMD and Multi-value IR Showcased for Static Verification of Collectives”. In: *Recent Advances in the Message Passing Interface*. Cham: Springer Nature Switzerland, 2025, pp. 3–20. ISBN: 978-3-031-73370-3. DOI: 10.1007/978-3-031-73370-3\_1.
- [3] OpenSHMEM Committee. *OpenSHMEM: Application Programming Interface Version 1.5*. 2020. URL: [http://openshmem.org/site/sites/default/site\\_files/OpenSHMEM-1.5.pdf](http://openshmem.org/site/sites/default/site_files/OpenSHMEM-1.5.pdf) (visited on 12/30/2024).
- [4] Timur Doumler and Jens Maurer. *A Natural Syntax for Contracts*. Nov. 8, 2023. URL: <https://www.open-std.org/jtc1/sc22/wg21/docs/papers/2023/p2961r2.pdf>.
- [5] Alexander Droste, Michael Kuhn, and Thomas Ludwig. “MPI-checker: Static Analysis for MPI”. In: *Proceedings of the Second Workshop on the LLVM Compiler Infrastructure in HPC*. LLVM ’15. New York, NY, USA: Association for Computing Machinery, Nov. 15, 2015, pp. 1–10. ISBN: 978-1-4503-4005-2. DOI: 10.1145/2833157.2833159. URL: <https://dl.acm.org/doi/10.1145/2833157.2833159> (visited on 11/04/2024).

## References II

---

- [6] GASPI Forum. *GASPI: Global Address Space Programming Interface 17.1*. 2017. URL: <https://raw.githubusercontent.com/GASPI-Forum/GASPI-Forum.github.io/master/standards/GASPI-17.1.pdf> (visited on 12/30/2024).
- [7] Message Passing Interface Forum. *MPI: A Message-Passing Interface Standard Version 4.1*. 2023. URL: <https://www.mpi-forum.org/docs/mpi-4.1/mpi41-report.pdf>.
- [8] Sayan Ghosh et al. “MiniVite: A Graph Analytics Benchmarking Tool for Massively Parallel Systems”. In: *2018 IEEE/ACM Performance Modeling, Benchmarking and Simulation of High Performance Computer Systems (PMBS)*. 2018 IEEE/ACM Performance Modeling, Benchmarking and Simulation of High Performance Computer Systems (PMBS). Nov. 2018, pp. 51–56. DOI: 10.1109/PMBS.2018.8641631. URL: <https://ieeexplore.ieee.org/abstract/document/8641631> (visited on 01/19/2025).
- [9] Tobias Hilbrich et al. “MUST: A Scalable Approach to Runtime Error Detection in MPI Programs”. In: *Tools for High Performance Computing 2009*. Berlin, Heidelberg: Springer, 2010, pp. 53–66. ISBN: 978-3-642-11261-4. DOI: 10.1007/978-3-642-11261-4\_5.

## References III

---

- [10] Tim Jammer et al. “MPI-BugBench: A Framework for Assessing MPI Correctness Tools”. In: *Recent Advances in the Message Passing Interface*. Cham: Springer Nature Switzerland, 2025, pp. 121–137. ISBN: 978-3-031-73370-3. DOI: 10.1007/978-3-031-73370-3\_8.
- [11] I Karlin, J Keasler, and J Neely. *LULESH 2.0 Updates and Changes*. LLNL-TR-641973, 1090032. July 22, 2013, LLNL-TR-641973, 1090032. DOI: 10.2172/1090032. URL: <https://www.osti.gov/servlets/purl/1090032/> (visited on 01/19/2025).
- [12] *Kotlin Contracts Proposal*. GitHub. Dec. 26, 2019. URL: <https://github.com/Kotlin/KEEP/blob/master/proposals/kotlin-contracts.md> (visited on 02/07/2025).
- [13] Patrick Ohly and Werner Krotz-Vogel. “Automated MPI Correctness Checking What If There Was a Magic Option?” In: *Proceedings of the 8th LCI International Conference on High-Performance Clustered Computing*. May 2007, pp. 19–25.
- [14] T. J. Parr and R. W. Quong. “ANTLR: A Predicated-LL(k) Parser Generator”. In: *Software: Practice and Experience* 25.7 (1995), pp. 789–810. ISSN: 1097-024X. DOI: 10.1002/spe.4380250705. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1002/spe.4380250705> (visited on 11/30/2024).

## References IV

---


- [15] Emmanuelle Saillard et al. “Static Local Concurrency Errors Detection in MPI-RMA Programs”. In: *2022 IEEE/ACM Sixth International Workshop on Software Correctness for HPC Applications (Correctness)*. 2022 IEEE/ACM Sixth International Workshop on Software Correctness for HPC Applications (Correctness). Nov. 2022, pp. 18–26. DOI: 10.1109/Correctness56720.2022.00008. URL: <https://ieeexplore.ieee.org/abstract/document/10027512> (visited on 11/04/2024).
- [16] Simon Schwitanski et al. “RMARaceBench: A Microbenchmark Suite to Evaluate Race Detection Tools for RMA Programs”. In: *Proceedings of the SC '23 Workshops of The International Conference on High Performance Computing, Network, Storage, and Analysis*. SC-W '23. New York, NY, USA: Association for Computing Machinery, Nov. 12, 2023, pp. 205–214. ISBN: 979-8-4007-0785-8. DOI: 10.1145/3624062.3624087. URL: <https://dl.acm.org/doi/10.1145/3624062.3624087> (visited on 12/08/2024).
- [17] Rob F. Van der Wijngaart and Timothy G. Mattson. “The Parallel Research Kernels”. In: *2014 IEEE High Performance Extreme Computing Conference (HPEC)*. 2014 IEEE High Performance Extreme Computing Conference (HPEC). Sept. 2014, pp. 1–6. DOI: 10.1109/HPEC.2014.7040972. URL: <https://ieeexplore.ieee.org/abstract/document/7040972> (visited on 01/19/2025).

### Common False-Positive Causes

```
1 if (rank == 0) {  
2     MPI_Irecv(buf, ..., MPI_REQUEST_NULL);  
3 }  
4 if (rank == 1) {  
5     MPI_Isend(buf, ..., &mpi_request_0);  
6 }  
7 MPI_Wait(&mpi_request_0, MPI_STATUS_IGNORE);
```

- Branch conditions not checked
  - Future Work: Evaluate Multi-Value analysis for improvement
- Path with both conditions evaluated to true checked
- $\Rightarrow$  Invalid data race report
- Duplicate error: Unmatched Wait, Missing Wait

Absolute Compilation Time Comparison



### Output from introduction example

```
1  ## Contract violation detected! ##
2  --> Function: MPI_Comm_rank
3  --> Contract: ...
4      --> Message: "Missing Initialization call"
5  ...
6      --> Message: "Missing Finalization call"
7  // Init-/Finalization errors left out for other functions
8
9  ## Contract violation detected! ##
10 --> Function: MPI_Type_contiguous
11 --> Contract: ...
12      --> Message: "Type not committed before use"
```